



C++20 Quick Syntax Reference

A Pocket Guide to the Language,
APIs, and Library

—
Fourth Edition

—
Mikael Olsson

apress®

C++20 Quick Syntax Reference

**A Pocket Guide to the
Language, APIs, and Library**

Fourth Edition

Mikael Olsson

Apress®

C++20 Quick Syntax Reference: A Pocket Guide to the Language, APIs, and Library

Mikael Olsson
Hammarland, Finland

ISBN-13 (pbk): 978-1-4842-5994-8
<https://doi.org/10.1007/978-1-4842-5995-5>

ISBN-13 (electronic): 978-1-4842-5995-5

Copyright © 2020 by Mikael Olsson

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed.

Trademarked names, logos, and images may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, logo, or image we use the names, logos, and images only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

The use in this publication of trade names, trademarks, service marks, and similar terms, even if they are not identified as such, is not to be taken as an expression of opinion as to whether or not they are subject to proprietary rights.

While the advice and information in this book are believed to be true and accurate at the date of publication, neither the authors nor the editors nor the publisher can accept any legal responsibility for any errors or omissions that may be made. The publisher makes no warranty, express or implied, with respect to the material contained herein.

Managing Director, Apress Media LLC: Welmoed Spahr
Acquisitions Editor: Steve Anglin
Development Editor: Matthew Moodie
Coordinating Editor: Mark Powers

Cover designed by eStudioCalamar

Cover image designed by Freepik (www.freepik.com)

Distributed to the book trade worldwide by Springer Science+Business Media New York, 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax (201) 348-4505, e-mail orders-ny@springer-sbm.com, or visit www.springeronline.com. Apress Media, LLC is a California LLC and the sole member (owner) is Springer Science + Business Media Finance Inc (SSBM Finance Inc). SSBM Finance Inc is a **Delaware** corporation.

For information on translations, please e-mail editorial@apress.com; for reprint, paperback, or audio rights, please email bookpermissions@springernature.com.

Apress titles may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Print and eBook Bulk Sales web page at <http://www.apress.com/bulk-sales>.

Any source code or other supplementary material referenced by the author in this book is available to readers on GitHub via the book's product page, located at www.apress.com/9781484259948. For more detailed information, please visit <http://www.apress.com/source-code>.

Printed on acid-free paper

Table of Contents

About the Authorxiii

About the Technical Reviewer xv

Introduction xvii

Chapter 1: Hello World 1

 Choosing an IDE 1

 Creating a Project 2

 Adding a Source File 2

 Selecting Language Standard 2

 Hello World 3

 Using the Standard Namespace 4

 IntelliSense 4

Chapter 2: Compile and Run 5

 Visual Studio Compilation 5

 Console Compilation 5

 Comments 6

Chapter 3: Variables 7

 Data Types 7

 Declaring Variables 8

 Assigning Variables 8

 Variable Scope 9

 Integer Types 10

TABLE OF CONTENTS

Signed and Unsigned Integers	11
Numeric Literals.....	12
Floating-Point Types.....	13
Literal Suffixes	14
Char Type	14
Bool Type.....	16
Chapter 4: Operators	17
Arithmetic Operators.....	17
Assignment Operators	18
Increment and Decrement Operators.....	18
Comparison Operators	19
Logical Operators	19
Bitwise Operators	20
Operator Precedence	20
Chapter 5: Pointers	23
Creating Pointers	23
Dereferencing Pointers	24
Pointing to a Pointer.....	25
Dynamic Allocation	25
Null Pointer	26
Chapter 6: References	29
Creating References.....	29
References and Pointers.....	30
Reference and Pointer Guideline.....	30
Rvalue Reference	30

Chapter 7: Arrays.....	33
Array Declaration and Allocation.....	33
Array Assignment.....	33
Multidimensional Arrays	34
Dynamic Arrays.....	35
Array Size.....	35
Vector.....	36
Chapter 8: Strings.....	39
String Combining	39
Escape Characters	40
String Compare	41
String Functions.....	41
String Encodings.....	42
String Formatting.....	43
Chapter 9: Conditionals	45
If Statement	45
Switch Statement	46
Ternary Operator	47
Initializers.....	47
Chapter 10: Loops.....	49
While Loop	49
Do-while Loop.....	50
For Loop.....	50
Break and Continue.....	52
Goto Statement.....	52

TABLE OF CONTENTS

Chapter 11: Functions53

 Defining Functions53

 Calling Functions53

 Function Parameters.....54

 Default Parameter Values.....54

 Function Overloading55

 Return Statement.....55

 Forward Declaration.....56

 Pass by Value57

 Pass by Reference58

 Pass by Address.....58

 Return by Value, Reference, or Address59

 Inline Functions.....60

 Auto and Decltype.....61

 Returning Multiple Values63

 Lambda Functions.....66

Chapter 12: Classes71

 Class Methods.....71

 Inline Methods72

 Object Creation72

 Accessing Object Members73

 Forward Declaration.....74

Chapter 13: Constructors75

 Constructor Overloading76

 This Keyword77

 Field Initialization.....77

TABLE OF CONTENTS

Default Constructor	78
Destructor	79
Special Member Functions	79
Object Initialization	80
Direct Initialization	81
Value Initialization	81
Copy Initialization	82
New Initialization	82
Aggregate Initialization	83
Uniform Initialization	83
Designated Initializers	85
Chapter 14: Inheritance	87
Upcasting	87
Downcasting	88
Constructor Inheritance	89
Multiple Inheritance	90
Chapter 15: Overriding	93
Hiding Derived Members	93
Overriding Derived Members	94
Base Class Scoping	96
Pure Virtual Functions	96
Chapter 16: Access Levels	99
Private Access	99
Protected Access	100
Public Access	101

TABLE OF CONTENTS

Access Level Guideline	101
Friend Classes and Functions	102
Public, Protected, and Private Inheritance	104
Chapter 17: Static	105
Static Fields	105
Static Methods	106
Static Local Variables	107
Static Global Variables	107
Chapter 18: Enum Types	109
Enum Example	109
Enum Constant Values	110
Enum Scope	111
Weakly Typed Enums	111
Enum Constant Type	113
Chapter 19: Structs and Unions	115
Structs	115
Struct Initialization	115
Union	117
Anonymous Union	118
Chapter 20: Operator Overloading	121
Binary Operator Overloading	122
Unary Operator Overloading	122
Comparison Operator Overloading	123
Overloadable Operators	126

Chapter 21: Custom Conversions	127
Implicit Conversion Constructor	127
Explicit Conversion Constructor	129
Conversion Operators.....	129
Explicit Conversion Operators	130
Chapter 22: Namespaces.....	133
Accessing Namespace Members.....	134
Nesting Namespaces	134
Importing Namespaces	135
Namespace Member Import	135
Namespace Alias.....	136
Type Alias	136
Including Namespace Members	137
Chapter 23: Constants	139
Constant Variables	139
Constant Pointers.....	139
Constant References.....	140
Constant Objects	140
Constant Methods	141
Constant Return Type and Parameters.....	142
Constant Fields	142
Constant Expressions.....	143
Immediate Functions	146
Constant Guideline	146

TABLE OF CONTENTS

Chapter 24: Preprocessor	147
Including Source Files.....	148
Define.....	148
Undefine.....	149
Predefined Macros	149
Macro Functions	150
Conditional Compilation	152
Compile if Defined.....	153
Error	153
Line	154
Pragma.....	154
Attributes	154
Chapter 25: Exception Handling	157
Throwing Exceptions.....	157
Try-Catch Statement	157
Rethrowing Exceptions	159
Noexcept Specifier.....	159
Exception Class.....	160
Chapter 26: Type Conversions	161
Implicit Conversions.....	161
Explicit Conversions.....	162
C++ Casts	162
Static Cast.....	163
Reinterpret Cast	163

Const Cast.....	164
C-Style and New-Style Casts	165
Dynamic Cast	165
Dynamic or Static Cast.....	167
Chapter 27: Smart Pointers	169
Unique Pointer	169
Shared Pointer	170
Weak Shared Pointer	171
Chapter 28: Templates	173
Function Templates	173
Calling Function Templates	174
Multiple Template Parameters	175
Class Templates	176
Non-type Parameters	177
Default Types and Values	178
Class Template Specialization.....	178
Function Template Specialization	179
Variable Templates.....	180
Variadic Templates	181
Fold Expressions	182
Concepts	184
Abbreviated Function Templates.....	186
Template Lambdas.....	187

TABLE OF CONTENTS

Chapter 29: Headers189

 Why Use Headers 189

 Using Headers..... 190

 What to Include in Headers 190

 Inline Variables..... 194

 Include Guards 195

 Modules 196

Index.....199

About the Author

Mikael Olsson is a professional web entrepreneur, programmer, and author. He works for an R&D company in Finland where he specializes in software development.

In his spare time, he writes books and creates websites that summarize various fields of interest. The books he writes are focused on teaching their subject in the most efficient way possible, by explaining only what is relevant and practical without any unnecessary repetition or theory.

About the Technical Reviewer

Marc Gregoire is a software engineer from Belgium. He graduated from the University of Leuven, Belgium, with a degree in “Burgerlijk ingenieur in de computerwetenschappen” (equivalent to a master of science degree in computer engineering). The year after, he received the *cum laude* degree of master in artificial intelligence at the same university. After his studies, Marc started working for a software consultancy company called Ordina Belgium. As a consultant, he worked for Siemens and Nokia Siemens Networks on critical 2G and 3G software running on Solaris for telecom operators. This required working on international teams stretching from South America and the United States to Europe, the Middle East, and Asia. Currently, Marc works for Nikon Metrology on industrial 3D laser scanning software.

Introduction

The C++ programming language is a general-purpose multiparadigm language created by Bjarne Stroustrup. The development of the language started in 1979 under the name “C with classes.” As the name implies, it was an extension of the C language with the additional concept of classes. Stroustrup wanted to create a better C that combined the power and efficiency of C with high-level abstractions to better manage large development projects. The resulting language was renamed C++ (pronounced “C-plus-plus”) in 1983. As a deliberate design feature, C++ maintains compatibility with C, and so most C code can easily be made to compile in C++.

The introduction of C++ became a major milestone in the software industry as a widely successful language for both system and application development. System programming involves software that controls the computer hardware directly, such as drivers, operating systems, and software for embedded microprocessors. These areas remain the core domain of the language, where resources are scarce and come at a premium. C++ is also widely used for writing applications, which run on top of system software, especially high-performance software such as games, databases, and resource-demanding desktop applications. Despite the introduction of many modern, high-level languages in this domain—such as Java, C#, and Python—C++ still holds its own and overall remains one of the most popular and influential programming languages in use today.

There are several reasons for the widespread adoption of C++. The foremost reason was the rare combination of high-level and low-level abstractions from the hardware. The low-level efficiency was inherited

INTRODUCTION

from C, and the high-level constructs came in part from a simulation language called Simula. This combination makes it possible to write C++ software with the strength of both approaches. Another strong point of the language is that it does not impose a specific programming paradigm on its users. It is designed to give the programmer a lot of freedom by supporting many different programming styles or paradigms, such as procedural, object-oriented, and generic programming.

C++ is updated and maintained by the C++ standards committee. In 1998, the first international standard was published, known informally as C++98. The language has since undergone five more revisions with further improvements, including C++03, C++11, C++14, C++17, and most recently C++20, which is the latest ISO standard for the C++ programming language released in 2020.

CHAPTER 1

Hello World

Choosing an IDE

To begin developing in C++, you need a text editor and a C++ compiler. You can get both at the same time by installing an Integrated Development Environment (IDE) that includes support for C++. A good choice is Microsoft's Visual Studio Community Edition, which is a free version of Visual Studio that is available from Microsoft's website.¹ The C++ compiler that comes with this IDE has good support for the C++17 standard and includes many features of C++20 as of the 2019 version. If you are running the Visual Studio installer on Windows, make sure to select the “Desktop development with C++” workload to enable development in C++.

Visual Studio is available on Windows and Mac, and there is a lightweight version called Visual Studio Code which can also be run on Linux. Two other popular cross-platform IDEs include NetBeans and Eclipse CDT. Alternatively, you can develop using a simple text editor such as Notepad, although this is less convenient than using an IDE. If you choose to use a simple text editor, just create an empty document with a .cpp file extension and open it in the editor of your choice.

¹<http://visualstudio.microsoft.com>

Creating a Project

After installing Visual Studio 2019, go ahead and launch the program. You then need to create a project, which will manage the C++ source files and other resources. Go to **File ► New ► Project** in Visual Studio to display the Create a new project window. From there, select the C++ language from the drop-down list to view only the C++ project templates. Then select the Empty Project template and click the Next button. At the next screen, you can configure the name and location of the project if you want to. When you are finished, click the Create button to let the wizard create your empty project.

Adding a Source File

You have now created a C++ project. In the Solution Explorer pane (choose **View ► Solution Explorer**), you can see that the project consists of three empty folders: Header Files, Resource Files, and Source Files. Right-click the Source Files folder and choose **Add ► New Item**. From the Add New Item dialog box, choose the C++ File (.cpp) type. Give this source file the name `MyApp` and click the Add button. An empty .cpp file will now be added to your project and opened for you.

Selecting Language Standard

To enable the latest features of the C++ language outlined in this book, it is necessary to manually change the language standard setting for your project. You can do this by first going to **Project ► Properties** to bring up the Property pages. From there, navigate to **Configuration Properties ► C/C++ ► Language ► C++ Language Standard**. Select the latest standard from the drop-down list (`std:c++latest`). Click OK and the project will now be configured to compile with the latest supported C++20 features.

Hello World

The first thing to add to the source file is the `main()` function. This is the entry point of the program, and the code inside of the curly brackets is executed when the program runs. The brackets, along with their content, are referred to as a code block, or just a block.

```
int main() {}
```

The first application will simply output the text "Hello World" to the screen. Before this can be done, the `iostream` header needs to be included. This header provides input and output functionality for the program, and it is one of the standard library files that comes with all C++ compilers. The `#include` directive effectively replaces the line with everything in the specified header before the file is compiled into an executable.

```
#include <iostream>
int main() {}
```

With `iostream` included, you gain access to several new functions. These are all located in the standard namespace called `std`, which you can examine by using a double colon, also called the *scope resolution operator* (`::`). After typing this in Visual Studio, the IntelliSense window will automatically open, displaying the namespace contents. Among the members, you find the `cout` stream. This is the standard output stream in C++ which can be used to print text to a console window. It uses two less than signs, collectively known as the insertion operator (`<<`), to indicate what to output. The string can then be specified, delimited by double quotes, and followed by a semicolon. The semicolon is used in C++ to mark the end of a statement.

```
#include <iostream>

int main()
{
    std::cout << "Hello World";
}
```

Using the Standard Namespace

To make things a bit easier, you can add a `using` directive to specify that this code file uses the standard namespace. You then no longer have to prefix `cout` with the namespace (`std::`) since it is used by default.

```
#include <iostream>
using namespace std;

int main()
{
    cout << "Hello World";
}
```

IntelliSense

When writing code in Visual Studio, a window called *IntelliSense* will pop up wherever there are multiple predetermined alternatives from which to choose. This window can also be brought up manually at any time by pressing `Ctrl+Space` to provide quick access to any code entities you are able to use within your program. This is a very powerful feature that you should learn to make good use of.

CHAPTER 2

Compile and Run

Visual Studio Compilation

Continuing from the last chapter, the Hello World program is now complete and ready to be compiled and run. You can do this by going to the Debug menu and clicking Start Without Debugging (Ctrl+F5). Visual Studio then compiles and runs the application, which displays the text in a console window.

Console Compilation

As an alternative to using an IDE, you can also compile source files from a terminal window as long as you have a C++ compiler.¹ For example, on a Linux machine, you can use the GNU C++ compiler, which is available on virtually all UNIX systems, including Linux and the BSD family, as part of the GNU Compiler Collection (GCC). This compiler can also be installed on Windows by downloading MinGW or on the Mac as part of the Xcode development environment.

¹www.stroustrup.com/compilers.html

To use the GNU compiler, you type its name `g++` in a terminal window and give it the input and output file names as arguments. It then produces an executable file, which when run gives the same result as the one compiled in Visual Studio.

```
g++ MyApp.cpp -o MyApp
./MyApp
Hello World
```

Comments

Comments are used to insert notes into the source code. They have no effect on the end program and are meant only to enhance the readability of the code, both for you and for other developers. C++ has two kinds of comment notations: single-line and multiline. The single-line comment starts with `//` and extends to the end of the line.

```
// single-line comment
```

The multiline comment may span more than one line and is delimited by `/*` and `*/`.

```
/* multi-line
comment */
```

Keep in mind that whitespace characters—such as spaces and tabs—are generally ignored by the compiler. This gives you a lot of freedom in how to format your code.

CHAPTER 3

Variables

Variables are used for storing data in memory during program execution.

Data Types

Depending on the type of data you need to store, there are several kinds of built-in data types. These are often called fundamental data types or *primitives*. The integer (whole number) types are short, int, long, and long long. The float, double, and long double types are floating-point (real number) types. The char type holds a single character, and the bool type contains either a true or false value.

Data Type	Size (Byte)	Description
char	1	Integer or character
short	2	Integer
int	4	Integer
Long	4 or 8	Integer
long long	8	Integer
float	4	Single-precision floating number
double	8	Double-precision floating number
long double	8 or 16	Floating-point number
bool	1	Boolean value

In C++, the exact size and range of primitive data types are not defined by the standard. Instead, they are dependent on the system for which the program is compiled. The sizes shown in the previous table are found on most 32-bit systems and are given in C++ bytes. A byte in C++ is the minimum addressable unit of memory which is guaranteed to be at least 8 bits, but might also be 16 or 32 bits depending on the system. By definition, a `char` in C++ is 1 byte in size. Furthermore, the `int` type will be 32 bits in size on 32-bit and 64-bit systems. Each integer type in the table must be at least as large as the one preceding it. The same applies to floating-point types, where each one must provide at least as much precision as the preceding one.

Declaring Variables

To *declare* (create) a variable, you start with the data type you want the variable to hold followed by an *identifier*, which is the name of the variable. The name can consist of letters, numbers, and underscores, but it cannot start with a number. It also cannot contain spaces or special characters and must not be a reserved keyword.

```
int myInt;    // correct
int 32Int;    // incorrect (starts with number)
int Int 32;   // incorrect (contains space)
int Int@32;   // incorrect (contains special character)
int new;      // incorrect (reserved keyword)
```

Assigning Variables

To assign a value to a declared variable, you use an equals sign, which is called the assignment operator (=).

```
myInt = 50;
```

The declaration and assignment can be combined into a single statement. When a variable is assigned a value, it then becomes *defined*.

```
int myInt = 50;
```

At the same time that the variable is declared, there are two alternative ways of assigning, or *initializing*, it by enclosing the value in either parentheses or braces. These examples are equivalent to the previous statement.

```
int myInt2(50); // direct initialization
int myInt3{50}; // uniform initialization
```

If you need to create more than one variable of the same type, there is a shorthand way of doing this using the comma operator (,).

```
int x = 1, y = 2, z;
```

Once a variable has been defined (declared and assigned), you can use it by simply referencing the variable's name, for example, to print it. Note the use of the `endl` stream manipulator token here to add a line break to the output stream.

```
cout << x << y << endl; // "12"
```

Variable Scope

The scope of a variable refers to the region of code within which it is possible to use that variable. Variables in C++ may be declared both globally and locally. A global variable is declared outside of any code blocks and is accessible from anywhere after it has been declared. A local variable, on the other hand, is declared inside of a function and will only be accessible within that function after it has been declared. The lifetime of a local variable is also limited. A global variable will remain allocated for the duration of the program, while a local variable will be destroyed when its function has finished executing.

CHAPTER 3 VARIABLES

```
int globalVar; // global variable
int main() { int localVar; } // local variable
```

The default values for these variables are also different. Global variables are automatically initialized to zero by the compiler, whereas local variables are not initialized at all. Uninitialized local variables will therefore contain whatever garbage is already present in that memory location.

```
int globalVar; // initialized to 0

int main()
{
    int localVar; // uninitialized
}
```

Using uninitialized variables is a common programming mistake that can produce unexpected results. It is therefore a good idea to always give your local variables an initial value when they are declared.

```
int main()
{
    int localVar = 0; // initialized to 0
}
```

Integer Types

There are four integer types you can use depending on how large a number you need the variable to hold.

```
char myChar = 0;    // -128 to +127
short myShort = 0;  // -32768 to +32767
int myInt = 0;       // -231 to +231-1
long myLong = 0;     // -231 to +231-1
```

C++11 standardized a fifth integer type, `long long`, which is guaranteed to be at least 64 bits large. Many compilers started to support this data type well before the C++11 standard was complete, including the Microsoft C++ compiler.

```
long long myL2 = 0; // -2^63 to +2^63-1
```

To determine the exact size of a data type, you can use the `sizeof` operator. This operator returns the number of bytes that a data type occupies in the system you are compiling for.

```
cout << sizeof(myChar) // 1 byte (per definition)
      << sizeof(myShort) // 2
      << sizeof(myInt) // 4
      << sizeof(myLong) // 4
      << sizeof(myL2); // 8
```

Fixed-sized integer types were added in C++11. These types belong to the `std` namespace and can be included through the `cstdint` standard library header.

```
#include <stdint>
using namespace std;

int8_t myInt8 = 0; // 8 bits
int16_t myInt16 = 0; // 16 bits
int32_t myInt32 = 0; // 32 bits
int64_t myInt64 = 0; // 64 bits
```

Signed and Unsigned Integers

By default, all integer types are signed and may therefore contain both positive and negative values. To explicitly declare a variable as signed, use the `signed` keyword.

CHAPTER 3 VARIABLES

```
signed char myChar = 0; // -128 to +127
signed short myShort = 0; // -32768 to +32767
signed int myInt = 0; // -2^31 to +2^31-1
signed long myLong = 0; // -2^31 to +2^31-1
signed long long myL2 = 0; // -2^63 to +2^63-1
```

If you only need to store positive values, you can declare integer types as unsigned to double their upper range.

```
unsigned char myChar = 0; // 0 to 255
unsigned short myShort = 0; // 0 to 65535
unsigned int myInt = 0; // 0 to 2^32-1
unsigned long myLong = 0; // 0 to 2^32-1
unsigned long long myL2 = 0; // 0 to 2^64-1
```

The signed and unsigned keywords may be used as stand-alone types, which are short for signed int and unsigned int.

```
unsigned uInt; // unsigned int
signed sInt; // signed int
```

Similarly, the short and long data types are abbreviations of short int and long int.

```
short myShort; // short int
long myLong; // long int
```

Numeric Literals

In addition to standard decimal notation, integers can also be assigned by using octal or hexadecimal notation. Octal literals use the prefix 0 and hexadecimal literals start with 0x. Both numbers shown here represent the same number, which in decimal notation is 50.

```
int myOct = 062; // octal notation (0)
int myHex = 0x32; // hexadecimal notation (0x)
```

As of C++14, there is a binary notation as well, which uses `0b` as its prefix. This version of the standard also added a digit separator (`'`), which can make it easier to read long numbers. The following binary number represents 50 in decimal notation.

```
int myBin = 0b0011'0010; // binary notation (0b)
```

Floating-Point Types

The floating-point types can store real numbers with different levels of precision.

```
float myFloat;    // ~7 digits
double myDouble;  // ~15 digits
long double myLD; // typically same as double
```

The precision shown here refers to the total number of digits in the number. A float can accurately represent about 7 digits, whereas a double can handle around 15. Trying to assign more than seven digits to a float means that the least significant digits will get rounded off.

```
myFloat = 12345.678; // rounded to 12345.68
```

Floats and doubles can be assigned by using either decimal or exponential notation. Exponential (scientific) notation is used by adding `E` or `e` followed by the decimal exponent.

```
myFloat = 3e2; // 3*10^2 = 300
```

As of C++17, the base may be specified as a hexadecimal value using the `0x` prefix. For such a number, the exponent part may use `p` instead of `e` to have the significant be scaled to the power of 2 rather than 10.

```
myFloat = 0xFp2; // 15*2^2 = 60
```

Literal Suffixes

An integer literal (constant) is normally treated as an `int` by the compiler, or a larger type if needed to fit the value. Suffixes can be added to the literal to change this evaluation. With integers, the suffix can be a combination of `U` and `L`, for unsigned and long, respectively. C++11 also added the `LL` suffix for the long long type. The order and casing of these letters do not matter.

```
int i = 10;  
long l = 10L;  
unsigned long ul = 10UL;
```

A floating-point literal is treated as a `double` unless otherwise specified. The `F` or `f` suffix can be used to specify that a literal is of the `float` type instead. Likewise, the `L` or `l` suffix specifies the long double type.

```
float f = 1.23F;  
double d = 1.23;  
long double ld = 1.23L;
```

The compiler implicitly converts literals to whichever type is necessary, so this type distinction for literals is usually not necessary. If the `F` suffix is left out when assigning to a `float` variable, the compiler may give a warning since the conversion from `double` to `float` involves a loss of precision.

Char Type

The `char` type is commonly used to represent ASCII characters. Such character constants are enclosed in single quotes and can be stored in a variable of `char` type.

```
char c = 'x'; // assigns 120 (ASCII for 'x')
```

The conversion between the number stored in the `char` and the character shown when the `char` is printed occurs automatically.

```
cout << c; // prints 'x'
```

For another integer type to be displayed as a character, it has to be explicitly cast to `char`. The recommended way of doing this is to use a `static_cast` as illustrated in the following, where the desired type is placed within angle brackets. Another way to perform the type cast is to use the legacy C-style cast, by placing the desired data type in parentheses before the variable or constant that is to be converted.

```
int i = c; // assigns 120
cout << i; // prints 120

// Prints 'x'
cout << static_cast<char>(i); // C++ new-style cast
cout << (char)i; // C-style cast
```

There are many ways to represent a character. Typically, ASCII encoding is used by most C++ compilers. In cases where code portability is important, this encoding can be assured by placing a `u8` prefix before the `char` literal. This prefix was added in C++17 and denotes the UTF-8 (Unicode) encoding, of which ASCII is a subset (the first 128 characters).

```
char ascii = u8'x'; // use UTF-8 encoding
```

UTF-16 and UTF-32 encodings can be represented using the `char16_t` and `char32_t` types, respectively, which were added in C++11. For completeness, C++20 added the `char8_t` type as well, which is used to represent a UTF-8 character and behaves the same as an unsigned `char`. The prefix `U` denotes a UTF-32 character and the `u` prefix a UTF-16 character.


```
char8_t c8 = 'A'; // UTF-8 character
char16_t c16 = u'€'; // UTF-16 character
char32_t c32 = U'𐀀'; // UTF-32 character
```

Bool Type

The `bool` type can store a Boolean value, which is a value that can only be either true or false. These values are specified with the `true` and `false` keywords.

```
bool b = false;
```

When used in an integer context, the Boolean value `false` is converted to zero and `true` is converted to one. Conversely, any value other than zero will be evaluated as `true` in a Boolean context. Note that the following `int` to `bool` conversion is made explicit, as the truncation would otherwise give a compiler warning.

```
int i = false; // 0
int j = true; // 1
bool b = static_cast<bool>(32); // true
```

CHAPTER 4

Operators

Operators are special symbols used to operate on values. The operators that deal specifically with numbers can be grouped into five types: arithmetic, assignment, comparison, logical, and bitwise operators.

Arithmetic Operators

The arithmetic operators include the four basic arithmetic operations, as well as the modulus operator (%), which is used to obtain the division remainder.

```
int i = 3 + 2; // 5, addition
    i = 3 - 2; // 1, subtraction
    i = 3 * 2; // 6, multiplication
    i = 3 / 2; // 1, division
    i = 3 % 2; // 1, modulus (division remainder)
```

Notice that the division operator gives an incorrect result. This is because it operates on two integer values and will therefore truncate the result and return an integer. To get the correct value, one of the numbers must be explicitly converted to a floating-point number in one of the following ways.

```
float f1 = 3 / 2.0f; // specify as floating-point number
float f2 = 3 / static_cast<float>(2); // C++ new-style cast
float f3 = 3 / (float)2; // C-style cast
```

Assignment Operators

The next group is the assignment operators. Most important is the assignment operator (=) itself, which assigns a value to a variable.

```
int i = 0; // assignment
```

A common use of the assignment and arithmetic operators is to operate on a variable and then save the result back into that same variable. These operations can be shortened with the combined assignment operators.

```
i += 5; // i = i+5;
i -= 5; // i = i-5;
i *= 5; // i = i*5;
i /= 5; // i = i/5;
i %= 5; // i = i%5;
```

Increment and Decrement Operators

Another common operation is to increment or decrement a variable by one. This can be simplified with the increment (++) and decrement (--) operators.

```
i++; // i = i+1;
i--; // i = i-1;
```

Both of these can be used either before or after a variable.

```
i++; // post-increment
i--; // post-decrement
++i; // pre-increment
--i; // pre-decrement
```

The result on the variable is the same whichever is used. The difference is that the post operator returns the original value before it changes the variable, while the pre operator changes the variable first and then returns the value.

```
int x, y;  
x = 5; y = x++; // y=5, x=6  
x = 5; y = ++x; // y=6, x=6
```

Comparison Operators

The comparison operators compare two values and return true or false. They are mainly used to specify conditions, which are expressions that evaluate to true or false.

```
bool b = (2 == 3); // equal to (false)  
b = (2 != 3); // not equal to (true)  
b = (2 > 3); // greater than (false)  
b = (2 < 3); // less than (true)  
b = (2 >= 3); // greater than or equal to (false)  
b = (2 <= 3); // less than or equal to (true)
```

Logical Operators

The logical operators are often used together with the comparison operators. “Logical and” (&&) evaluates to true if both the left and right sides are true, and “logical or” (||) is true if either the left or right side is true. For inverting a Boolean result, there is the logical not (!) operator. Note that for both “logical and” and “logical or,” the right side will not be evaluated if the result is already determined by the left side. This behavior is called short-circuiting.

```
bool b = (true && false); // logical and (false)
b = (true || false);      // logical or (true)
    b = !(true);          // logical not (false)
```

Bitwise Operators

The bitwise operators can manipulate individual bits inside an integer. For example, the “bitwise or” operator (`|`) makes the resulting bit 1 if the bits are set on either side of the operator.

```
int i = 5 & 4; // 101 & 100 = 100 (4)      // and
i = 5 | 4;     // 101 | 100 = 101 (5)      // or
i = 5 ^ 4;     // 101 ^ 100 = 001 (1)      // xor
i = 4 << 1;    // 100 << 1 = 1000 (8)      // left shift
i = 4 >> 1;    // 100 >> 1 = 10 (2)        // right shift
i = ~4;       // ~00000100 = 11111011 (-5) // invert
```

The bitwise operators also have combined assignment operators.

```
int i=5; i &= 4; // 101 & 100 = 100 (4) // and
    i=5; i |= 4; // 101 | 100 = 101 (5) // or
    i=5; i ^= 4; // 101 ^ 100 = 001 (1) // xor
    i=5; i <<= 1; // 101 << 1 = 1010 (10) // left shift
    i=5; i >>= 1; // 101 >> 1 = 10 (2) // right shift
```

Operator Precedence

In C++, expressions are normally evaluated from left to right. However, when an expression contains multiple operators, the precedence of those operators decides the order in which they are evaluated. The order of precedence can be seen in the following table, where the operator with the lowest precedence will be evaluated first. This same basic order also applies to many other languages, such as C, Java, and C#.

Pre	Operator	Pre	Operator
1	::	10	== !=
2	() [] . -> x++ x--	11	&
3	! ~ ++x --x *x &x (type) sizeof co_await new new[] delete delete[]	12	^
4	.* ->*	13	
5	* / %	14	&&
6	+ -	15	
7	<< >>	16	?: = op= throw co_yield
8	<=>	17	,
9	< <= > >=		

To give an example, “logical and” (&&) binds weaker than relational operators, which in turn bind weaker than arithmetic operators.

```
bool b = 2+3 > 1*4 && 5/5 == 1; // true
```

To make things clearer, parentheses can be used to specify which part of the expression will be evaluated first. As seen in the table, parentheses are among the operators with the greatest precedence.

```
bool b = ((2+3) > (1*4)) && ((5/5) == 1); // true
```

CHAPTER 5

Pointers

A *pointer* is a variable that contains the memory address of another variable, function, or object, called the *pointee*.

Creating Pointers

Pointers are declared as any other variable, except that an asterisk (*) is placed between the data type and the pointer's name. The data type used determines what type of memory it will point to. More than one pointer can be created in the same statement using the comma operator. The asterisk must then be placed before each identifier and not after the type.

```
int* p; // pointer to an integer
int *q; // alternative syntax
int *a, *b, *c; // multiple pointers
```

A pointer can point to a variable of the same type by prefixing that variable with an ampersand, in order to retrieve its address and assign it to the pointer. The ampersand is known as the address-of operator (&).

```
int i = 10;
p = &i; // address of i assigned to p
```

Dereferencing Pointers

The pointer now contains the memory address to the integer variable. Referencing the pointer will retrieve this address. To obtain the actual value stored in that address, the pointer must be prefixed with an asterisk, known as the dereference operator (*).

```
#include <iostream>
using namespace std;

int main()
{
    int i = 10;
    int* p = &i;
    cout << "Address of i: " << p << endl; // ex. 0017FF1C
    cout << "Value of i: " << *p << endl; // 10
}
```

When writing to the pointer, the same method is used. Without the asterisk, the pointer is assigned a new memory address, and with the asterisk the actual value of the variable pointed to will be updated.

```
p = &i; // address of i assigned to p
*p = 20; // value of i changed through p
```

If a second pointer is created and assigned the value of the first pointer, it will then get a copy of the first pointer's memory address.

```
int* p2 = p; // copy of p (copies address stored in p)
```


Pointing to a Pointer

Sometimes it can be useful to have a pointer that can point to another pointer. This is done by declaring a pointer with two asterisks and then assigning it the address of the pointer that it will reference. This way, when the address stored in the first pointer changes, the second pointer can follow that change.

```
int** r = &p; // pointer to p (assigns address of p)
```

Referencing the second pointer now gives the address of the first pointer. Dereferencing the second pointer gives the address of the variable, and dereferencing it again gives the value of the variable.

```
cout << "Address of p: " << r << endl; // ex. 0017FF28
cout << "Address of i: " << *r << endl; // ex. 0017FF1C
cout << "Value of i: " << **r << endl; // 20
```

Dynamic Allocation

One of the main usages of pointers is to allocate memory during runtime—so-called *dynamic allocation*. In the examples so far, the programs have only had as much memory available as was declared for the variables at compile time. This is referred to as *static allocation*, and those variables are stored on the so-called *stack*. If any additional memory is needed at runtime, the new operator has to be used. This operator allows for dynamic allocation of memory, which can only be accessed through pointers and is stored on the so-called *heap*. The new operator takes either a primitive data type or an object type as its argument, and it will return a pointer to the allocated memory as long as there is sufficient memory available.

```
int* d = new int; // dynamic allocation
```

An important thing to know about dynamic allocation is that the allocated memory will not be released like the rest of the program memory when it is no longer required. Instead, it has to be manually released with the `delete` keyword. This allows you to control the lifetime of a dynamically allocated object, but it also means that you are responsible for deleting it once it is no longer needed. Forgetting to delete memory that has been allocated with the `new` keyword will give the program memory leaks, because that memory will stay allocated until the program shuts down.

```
delete d; // release allocated memory
```

In modern C++, the use of the so-called smart pointers is preferred over regular pointers as they remove the need for manually deleting dynamically allocated memory. These pointers will be covered in a later chapter.

Null Pointer

A pointer should be set to null when it is not assigned to a valid address. Such a pointer is called a *null pointer*. Doing this will allow you to check whether the pointer can be safely dereferenced, because a valid pointer will never be null. In the early days before C++11, the constant `NULL` or the integer zero was used to symbolize the null pointer. The `NULL` constant is defined in the `cstdio` standard library file, which is included through `iostream`.

```
int* g = 0; // null pointer (unused pointer)
int* h = NULL; // null pointer
```

C++11 introduced the now preferred keyword `nullptr` to specify a null pointer, in order to distinguish between zero and a null pointer. The advantage of using `nullptr` is that unlike an integer zero, `nullptr`

will not implicitly convert to an integer type. The literal has its own type, `std::nullptr_t`, which can only be implicitly converted to pointer and bool types.

```
#include <iostream> // include nullptr_t type

int main()
{
    int* p = nullptr; // ok
    int i = nullptr; // error
    bool b = (bool) nullptr; // false
    std::nullptr_t mynull = nullptr; // ok
}
```

As seen earlier, a dynamically allocated object is accessed through a pointer and can be unallocated with the `delete` keyword. A point to keep in mind is that after deletion, the pointer will point to a now inaccessible memory location. Trying to dereference such a pointer will cause a runtime error.

```
int* m = new int; // allocate memory for object
delete m; // deallocate memory
*m = 5; // error: write access violation
```

To help prevent this, the deleted pointer should be set to null. Note that trying to delete an already deleted null pointer is safe. However, if the pointer has not been set to null, attempting to delete it again will cause memory corruption and possibly crash the program.

```
delete m;
m = nullptr; // mark as null pointer
delete m; // safe
```

CHAPTER 5 POINTERS

Since you may not always know whether a pointer is valid, a check should be made whenever a pointer is dereferenced to make sure that it is not null.

```
if (m != nullptr) { *m = 5; } // check for valid pointer
if (m) { *m = 5; } // alternative
```

CHAPTER 6

References

References allow programmers to create new names for a variable. They provide a simpler and safer alternative to pointers that should be used in favor of pointers whenever possible.

Creating References

A reference is declared in the same way as a regular variable, except that an ampersand is appended between the data type and the variable name. Furthermore, at the same time as the reference is declared, it must be initialized with a variable of the specified type.

```
int x = 5;
int& r = x; // r is an alias to x
int &s = x; // alternative syntax
int& t; // error: must be initialized
```

Once the reference has been assigned, or seated, it can never be reseated to another variable. The reference has in effect become an alias for the variable and can be used exactly as though it were the original variable.

```
r = 10; // assigns value to r/x
```

References and Pointers

A reference is similar to a pointer that always points to the same thing. However, while a pointer is a variable that points to another variable or object, a reference is only an alias and does not have an address of its own.

```
int* ptr = &r; // ptr assigned address to x
```

Reference and Pointer Guideline

Generally, whenever a pointer does not need to be reassigned, a reference should be used instead, because a reference is safer than a pointer since it must always refer to something. This means that there is no need to check if a reference refers to null, as should be done with pointers. It is possible for a reference to be invalid—for example, when a reference refers to a null pointer—but it is much easier to avoid this kind of mistake with references than it is with pointers.

```
int* ptr = nullptr; // null pointer
int& ref = *ptr;
ref = 10; // error: invalid memory access
```

Rvalue Reference

With C++11 came a new kind of reference called an rvalue reference. This reference can bind and modify temporary objects (rvalues), such as literal values and function return values. An rvalue reference is formed by placing two ampersands after the type.

```
int&& ref = 1 + 2; // rvalue reference
```

The rvalue reference extends the lifetime of the temporary object and allows it to be used like an ordinary variable.

```
ref += 3;  
cout << ref; // "6"
```

The benefit of rvalue references is that they allow unnecessary copying to be avoided when dealing with temporary objects. This offers greater performance, particularly when handling larger types, such as strings and objects.

CHAPTER 7

Arrays

An array is a data structure used for storing a collection of values that all have the same data type.

Array Declaration and Allocation

To declare an array, you start as you would a normal variable declaration, but in addition you append a set of square brackets following the array's name. The brackets contain the number of elements in the array.

```
int myArray[3]; // integer array with 3 elements
```

The default values for these elements are the same as for variables—elements in global arrays are initialized to their default values and elements in local arrays remain uninitialized.

Array Assignment

To assign values to the elements, you can reference them one at a time by placing the element's index inside the square brackets, starting with zero.

```
myArray[0] = 1;  
myArray[1] = 2;  
myArray[2] = 3;
```


Alternatively, you can assign values at the same time as the array is declared by enclosing the values in curly brackets. If the specified array length is larger than the number of values, the remaining values will be initialized to zero. The array length may optionally be left out to let the array size be decided by the number of values assigned.

```
int myArray[3] = { 1, 2, 3 };  
int myArray[] = { 1, 2, 3 };
```

Once the array elements are initialized, they can be accessed by referencing the index of the element you want.

```
int x = myArray[0]; // 1
```

Multidimensional Arrays

Arrays can be made multidimensional by adding more sets of square brackets. As with single-dimensional arrays, they can either be filled in one at a time or all at once during the declaration.

```
int myArray[2][2] = { { 0, 1 }, { 2, 3 } };  
myArray[0][0] = 0;  
myArray[0][1] = 1;
```

The extra curly brackets are optional, but including them is good practice since it makes the code easier to understand.

```
int myArray[2][2] = { 0, 1, 2, 3 }; // alternative
```

Dynamic Arrays

Because the previous arrays are made up of static (nondynamic) memory, their size must be determined before execution. Therefore, the size needs to be a constant value. In order to create an array with a size that is not known until runtime, you need to use dynamic memory, which is allocated with the `new` keyword and must be assigned to a pointer.

```
int* p = new int[3]; // dynamically allocated array
```

An array in C++ behaves as a constant pointer to the first element in the array. The referencing of array elements can therefore be made just as well with pointer arithmetic. By incrementing the pointer by one, you move to the next element in the array, because changes to a pointer's address are implicitly multiplied by the size of the pointer's data type.

```
*(p+1) = 10; // p[1] = 10;
```

Array Size

Just as with any other pointer, it is possible to exceed the valid range of an array and thereby rewrite some adjacent memory. This should always be avoided since it can lead to unexpected results or crash the program.

```
int myArray[2] = { 1, 2 };
myArray[2] = 3; // error: out of bounds
```

To determine the length of a regular (statically allocated) array, you can use the `std::size` function.

```
#include <iostream> // std::size
int main()
{
    int myArray[2] = { 1, 2 };
    int length = std::size(myArray); // 2
}
```

This method cannot be used for dynamically allocated arrays. The only way to determine the size of such an array is through the variable used in its allocation.

```
int size = 3;
int* p = new int[size]; // dynamically allocated array
```

When you are done using a dynamic array, you must remember to delete it. This is done using the `delete` keyword with an appended set of square brackets.

```
delete[] p; // release allocated array
p = nullptr; // mark pointer as unused
```

Vector

A vector is a container class representing a resizable array. The element type of the vector is specified in angle brackets after the class name, because vector is a so-called template class.

```
#include <vector> // std::vector
using namespace std;

int main()
{
    vector<int> myVector;
}
```

Vectors are preferable to dynamic arrays as they have a number of advantages including the ability to grow and shrink automatically as needed. Vectors will also implicitly deallocate themselves when they go out of scope, so there is no need to manually delete them. The following example illustrates how to assign, change, and read elements of a vector.

```
#include <vector>
using namespace std;

int main()
{
    vector<int> v;

    // Assign three elements with value two
    v.assign(3, 2); // [2, 2, 2]

    // Add 4 at last position
    v.push_back(4); // [2, 2, 2, 4]

    // Change first element
    v[0] = 1; // [1, 2, 2, 4]

    // Change second element (bound checked)
    v.at(2) = 3; // [1, 2, 3, 4]

    // Remove second element
    v.erase(v.begin()+1); // [1, 3, 4]

    // Remove last element
    v.pop_back(); // [1, 3]

    // Get vector length
    int len = v.size(); // 2

    // Print first and second elements
    cout << v.at(0) << v[1]; // "13"
}
```

CHAPTER 8

Strings

The `string` class in C++ is used to store string values. Before a string can be declared, the string header must be included. The standard namespace can also be used since the string class is part of that namespace.

```
#include <string>
using namespace std;
```

Strings can then be declared like any other data type. To assign a string value to a string variable, delimit the literals by double quotes and assign them to the variable. The initial value can also be assigned through either direct or uniform initialization at the same time as the string is declared.

```
string h = "Hello";
string w ("Hi"); // direct initialization
string u {"Hey"}; // uniform initialization
```

String Combining

The plus sign, known as the concatenation operator (+) in this context, is used to combine two strings. It has an accompanying assignment operator (+=) to append a string.

```
string a = "Hello";
string b = "World";
string c = a + " " + b; // Hello World
a += b; // HelloWorld
```

The concatenation operator will work as long as one of the strings it operates on is a C++ string. A string literal is by default a C-style string which consists of an array of char elements. The letter s can be appended to a string literal to instead represent it as the std::string type.

```
string d = "Hello" + "World"; // error, no C++ string
string e = "Hello" + "World"s; // ok
string f = e + "Again"; // ok
```

String literals will also be implicitly combined if the plus sign is left off.

```
string g = "Hel" "lo"; // ok
```

Escape Characters

A backslash notation is used to write special characters, such as the newline character \n.

```
string s = "Hello\nWorld";
```

These special characters are called escape characters, and they are described in the following table.

Character	Meaning	Character	Meaning
\n	Newline	\f	Form feed
\t	Horizontal tab	\a	Alert sound
\v	Vertical tab	\'	Single quote
\b	Backspace	\"	Double quote
\r	Carriage return	\\	Backslash
\0	Null character		

Additionally, any one of the 128 ASCII characters can be expressed by writing a backslash followed by the ASCII code for that character, represented as either an octal or hexadecimal number.

```
string oct = "\053"; // octal '+'  
string hex = "\x02B"; // hexadecimal '+'
```

As of C++11, escape characters can be ignored by adding an `R` before the string along with a set of parentheses within the double quotes. This is called a raw string literal and can be used, for instance, to make file paths more readable.

```
string escaped = "c:\\Windows\\System32\\cmd.exe";  
string raw = R"(c:\Windows\System32\cmd.exe)";
```

String Compare

The way to compare two strings is simply by using the equal to operator (`==`). This will not compare the memory addresses of the strings, as is the case with C strings.

```
string s = "Hello";  
bool b = (s == "Hello"); // true
```

String Functions

The string class has a lot of functions. Among the most useful ones are the length and size functions, which both return the number of characters in the string. Their return type is `size_t`, which is an unsigned data type used to hold the size of an object. This is simply an alias for one of the built-in

data types, but which one it is defined as varies between compilers. The alias is defined in the `cstdint` standard library header, which is included through `iostream`.

```
string s = "Hello";  
size_t i = s.length(); // 5, length of string  
i = s.size(); // 5, same as length()
```

Another useful function is `substr` (substring), which requires two parameters. The second parameter is the number of characters to return, starting from the position specified in the first parameter.

```
s.substr(0,2); // "He"
```

A single character from a string can also be extracted or changed by using the array notation.

```
char c = s[0]; // 'H'
```

String Encodings

A string enclosed within double quotes produces an array of the `char` type, which can only hold 256 unique symbols. To support larger character sets, the wide character type `wchar_t` is provided. Its size can vary between compilers so it is not platform independent. String literals of this type are created by prepending the string with a capital `L`. The resulting array can be stored using the `wstring` class. This class works like the basic `string` class but uses the `wchar_t` character type instead.

```
wstring s1 = L"Hello";  
wchar_t *s2 = L"Hello"; // C-style string
```


Fixed-size character types were introduced in C++11, namely, `char16_t` and `char32_t`. These types provide definite representations of the UTF-16 and UTF-32 encodings, respectively. UTF-16 string literals are prefixed with `u` and can be stored using the `u16string` class. Likewise, UTF-32 string literals are prefixed with `U` and are stored in the `u32string` class. The prefix `u8` was also added to represent a UTF-8 encoded string literal. A string consisting of UTF-8 literals can be stored in the `u8string` type added in C++20.

```
string s3 = "Compiler-defined encoding";
u8string s4 = u8"UTF-8 string";
u16string s5 = u"UTF-16 string";
u32string s6 = U"UTF-32 string";
```

Specific Unicode characters can be inserted into a string literal using the escape character `\u` followed by a hexadecimal number representing the character.

```
u8string s7 = u8"Asterisk: \u002A"; // "Asterisk: *"
```

String Formatting

C++20 introduced the `std::format` function as a more convenient and type-safe way to format strings compared with legacy string formatting functions such as the `printf` family inherited from C. The first argument to this function is the string to be formatted. Curly brackets (`{}`) appearing in the string will be replaced by successive arguments to the function as seen here.

```
// "1 plus 2 equals 3"
string f = std::format("1 plus 2 equals {}", 1+2);
```

CHAPTER 8 STRINGS

The curly brackets can include a number to specify which argument it will be replaced by.

```
// "5 is more than zero"  
string f = std::format("{1} is more than {0}", "zero", 5);
```

CHAPTER 9

Conditionals

Conditional statements are used to execute different code blocks based on different conditions.

If Statement

The `if` statement will execute only if the expression inside the parentheses is evaluated to true. In C++, this does not have to be a Boolean expression. It can be any expression that evaluates to a number, in which case zero is false and all other numbers are true.

```
if (x < 1) {  
    cout << x << " < 1";  
}
```

To test for other conditions, the `if` statement can be extended by any number of `else if` clauses.

```
else if (x > 1) {  
    cout << x << " > 1";  
}
```

The `if` statement can have one `else` clause at the end, which will execute if all previous conditions are false.

```
else {  
    cout << x << " == 1";  
}
```

As for the curly brackets, they can be left out if only a single statement needs to be executed conditionally. However, it is considered good practice to include them since they improve readability.

```
if (x < 1)  
    cout << x << " < 1";  
else if (x > 1)  
    cout << x << " > 1";  
else  
    cout << x << " == 1";
```

Switch Statement

The `switch` statement checks for equality between an integer and a series of case labels and then passes execution to the matching case. It may contain any number of case clauses as well as a default label for handling all other cases.

```
switch (x)  
{  
    case 0: cout << x << " is 0"; break;  
    case 1: cout << x << " is 1"; break;  
    default: cout << x << " is not 0 or 1"; break;  
}
```

Note that the statements after each case label end with the `break` keyword to skip the rest of the switch. If the `break` is left out, execution will fall through to the next case, which can be useful if several cases need to be evaluated in the same way.

Ternary Operator

In addition to the `if` and `switch` statements, there is the ternary operator (`?:`), which can replace a single `if/else` clause. This operator takes three expressions. If the first one is true, then the second expression is evaluated and returned, and if it is false, the third one is evaluated and returned.

```
x = (x < 0.5) ? 0 : 1; // ternary operator (?:)
```

C++ allows expressions to be used as stand-alone code statements. Because of this, the ternary operator cannot just be used as an expression, but also as a statement.

```
(x < 0.5) ? x = 0 : x = 1; // alternative syntax
```

The programming term *expression* refers to code that evaluates to a value, whereas a *statement* is a code segment that ends with a semicolon or a closing curly bracket.

Initializers

It is preferable to keep the scope of a variable limited to the section of code where the variable is used. This way, the variable is prevented from cluttering up the namespace unnecessarily or causing potential name clashes later on in the code. To assist with this, C++17 introduced the ability to declare and initialize a locally scoped variable for an `if`

statement, by adding an initializer before the condition. This reduces the scope of the variable so that it is only visible within the body of the `if` statement and any accompanying `else` clauses.

```
int a = 2, b = 3;
// ...
if (int sum = a+b; sum == 5) {
    cout << sum << " is 5";
}
```

Switch statements may also use an initializer as of C++17. Like the `if` statement, this feature helps avoid potential name clashes by limiting the scope of the variable to within the switch statement.

```
switch (int sum = a+b; sum) {
    case 5: cout << sum << " is 5"; break;
}
```

CHAPTER 10

Loops

There are three looping structures available in C++, all of which are used to execute a specific code block multiple times. Just as with the conditional `if` statement, the curly brackets for the loops can be left out if there is only one statement in the code block.

While Loop

The `while` loop runs through the code block only if its condition is true and will continue looping for as long as the condition remains true. Bear in mind that the condition is only checked at the start of each iteration (loop).

```
int i = 0;
while (i < 10) {
    cout << i++; // 0-9
}
```

Do-while Loop

The do-while loop works in the same way as the while loop, except that it checks the condition after the code block. It will therefore always run through the code block at least once, in contrast with the while loop. Notice that this loop ends with a semicolon.

```
int j = 0;
do {
    cout << j++; // 0-9
} while (j < 10);
```

For Loop

The for loop is used to run through a code block a set number of times. It uses three parameters. The first one initializes a counter and is always executed once before the loop. This counter variable is limited in scope to the for loop and is not accessible after the loop. The second parameter holds the condition for the loop and is checked before each iteration. Lastly, the third parameter contains the increment of the counter and is executed at the end of each loop.

```
for (int k = 0; k < 10; ++k) {
    cout << k; // 0-9
}
```

The for loop has several variations. For instance, the first and third parameters can be split into several statements by using the comma operator.

```
for (int k = 0, m = 0; k < 5; ++k, m--) {
    cout << k+m; // "00000"
}
```


There is also the option of leaving out any one of the parameters. The following are a couple of examples of this.

```
for (;;) {
    // infinite loop
}

for (int i=0; i<10; ) {
    // increment i inside of loop
}

int counter = 0;
for (; counter<10; ++counter) {
    // ...
}
// make counter usable outside of loop
```

C++11 introduced a range-based for loop syntax for iterating through arrays and other container types. At each iteration, the next element in the array is bound to the specified variable, in this case a reference variable, and the loop continues until it has gone through the entire array.

```
int a[3] = {1, 2, 3};
for (int &i : a) {
    cout << i; // "123"
}
```

C++20 extended the range-based for loop by allowing it to include an initializer. This is useful for keeping scopes tight when iterating over a temporary container that is only needed for the duration of the loop.

```
for (int a[3] = {1, 2, 3}; int &i : a) {
    cout << i; // "123"
}
```

Break and Continue

There are two jump statements that can be used inside loops: `break` and `continue`. The `break` keyword ends the loop structure, and `continue` skips the rest of the current iteration and continues at the beginning of the next iteration.

```
for (int i = 0; i < 10; i++)
{
    if (i == 5) break; // end loop
    if (i == 3) continue; // start next iteration
    cout << i; // "0124"
}
```

Goto Statement

A third jump statement that may be useful to know about is `goto`, which performs an unconditional jump to a specified label. This instruction is generally never used since it tends to make the flow of execution difficult to follow.

```
goto myLabel; // jump to label
myLabel: // label declaration
```

CHAPTER 11

Functions

Functions are reusable code blocks that will only execute when called.

Defining Functions

A function can be created by typing `void` followed by the function's name, a set of parentheses, and a code block. The `void` keyword means that the function will not return a value. A common naming convention for functions is to name them in the same way as variables—a descriptive name with each word initially capitalized, except for the first one.

```
#include <iostream>
using namespace std;

void myFunction()
{
    cout << "Hello World";
}
```

Calling Functions

The previous function will simply print out a text message when it is called. To invoke it from the main function, specify the function's name followed by a set of parentheses.

```
int main()
{
    myFunction(); // "Hello World"
}
```

Function Parameters

The parentheses that follow the function name are used to pass arguments to the function. To do this, you must first add the corresponding parameters to the function declaration in the form of a comma-separated list.

```
void myFunction(string a, string b)
{
    cout << a << " " << b;
}
```

A function can be defined to take any number of parameters, and those parameters can have any data types. Just ensure that the function is called with the same types and number of arguments.

```
myFunction("Hello", "World"); // "Hello World"
```

To be precise, *parameters* appear in function definitions, while *arguments* appear in function calls. However, the two terms are sometimes wrongly used.

Default Parameter Values

It is possible to specify default values for parameters by assigning them a value inside the parameter list.

```
void myFunction(string a, string b = "Earth")
{
    cout << a << " " << b;
}
```

Then, if that argument is unspecified when the function is called, the default value will be used instead. For this to work, it is important that the parameters with default values are to the right of those without default values.

```
myFunction("Hello"); // "Hello Earth"
```

Function Overloading

A function in C++ can be defined multiple times with different parameters. This is a powerful feature called *function overloading*, and it allows a function to handle a variety of arguments without the programmer needing to be aware of using different functions.

```
void myFunction(string a, string b) { cout << a << " " << b; }
void myFunction(string a) { cout << a; }
void myFunction(int a) { cout << a; }
```

Return Statement

A function can return a value. The `void` keyword is then replaced with the data type the function will return, and the `return` keyword must be added to the function's body followed by an argument of the specified return type. Keep in mind that all branches in the function must return a value.

```
int getSum(int a, int b)
{
    return a + b;
}
```

Return is a jump statement that causes the function to exit and return the specified value to the place where the function was called. For example, the previously defined function can be used with the output stream since the function evaluates to an integer.

```
cout << getSum(5, 10); // "15"
```

The return statement can also be used in void functions to exit before the end of the function block is reached.

```
void dummy() { return; }
```

Note that although the main function is set to return an integer type, it does not have to explicitly return a value. This is because the compiler will automatically add a `return 0` statement to the end of the main function.

```
int main() { return 0; }
```

Forward Declaration

An important point to keep in mind in C++ is that functions must be declared before they can be called. This does not mean that the function has to be implemented before it is called. It only means that the function's header needs to be specified at the beginning of the source file, so that the compiler knows that the function exists. This kind of forward declaration is known as a *prototype*.

```
void myFunction(int a); // prototype
int main()
{
    myFunction(0);
}
void myFunction(int a) {} // definition
```

The parameter names in the prototype do not need to be included. Only the data types must be specified. However, including the names serves as a kind of documentation, and they will also show up in IntelliSense, so it is a good practice to include them.

```
void myFunction(int);
```

Pass by Value

In C++, variables of both primitive and object data types are by default passed by value. This means that only a copy of the value or object is passed to the function. Therefore, changing the parameter in any way will not affect the original, and passing a large object will be slow.

```
#include <iostream>
#include <vector>
using namespace std;

void change(int i) { i = 10; }
void change(vector<int> a) { a.at(0) = 5; }

int main()
{
    int x = 0; // value type
    change(x); // copy of x is passed
    cout << x; // "0"

    vector<int> v { 3 }; // reference type
    change(v); // object copy is passed
    cout << v.at(0); // "3"
}
```

Pass by Reference

Alternatively, to instead pass a variable by reference, you just need to add an ampersand before the parameter's name in the function's definition. When arguments are passed by reference, both primitive and object data types can be changed, and the changes will affect the original variable.

```
void change(int& i) { i = 10; }

int main()
{
    int x = 0; // value type
    change(x); // reference is passed
    cout << x; // "10"
}
```

Pass by Address

As an alternative to passing by reference, arguments may also be passed by address using pointers. This passing technique serves the same purpose as passing by reference, but uses pointer syntax instead.

```
void change(int* i) { *i = 10; }

int main()
{
    int x = 0; // value type
    change(&x); // address is passed
    cout << x; // 10
}
```

One difference is that pointers can be null, whereas references cannot. So if the function should not allow null arguments, it is preferable to use pass by reference.

Return by Value, Reference, or Address

In addition to passing variables by value, reference, or address, a variable may also be returned in one of these ways. Normally, a function returns by value, in which case a copy of the value is returned to the caller.

```
int byVal(int i) { return i + 1; }

int main()
{
    int a = 10;
    cout << byVal(a); // "11"
}
```

To return by reference instead, an ampersand is placed after the function's return type. The function must then return a variable and may not return an expression or literal, as can be done when using return by value. The variable returned should never be a local variable, because the memory to these variables is released when the function ends. Instead, return by reference is commonly used to return an argument that has also been passed to the function by reference.

```
int& byRef(int& i) { return i; }

int main()
{
    int a = 10;
    cout << byRef(a); // "10"
}
```

To return by address, you append the dereference operator (*) to the function's return type. This return technique has the same two restrictions as when returning by reference—the address of a variable must be returned and the returned variable must not be local to the function.

```
int* byAdr(int* i) { return i; }

int main()
{
    int a = 10;
    cout << *byAdr(&a); // "10"
}
```

If a function returns a pointer, it may not be clear whether the function has allocated some memory dynamically and where this memory should be deallocated. For this reason, it is preferable to return by reference instead or to use smart pointers which are covered in a later chapter.

Inline Functions

A point to keep in mind when using functions is that every time a function is called, a small performance overhead occurs. To potentially remove this overhead, you can recommend that the compiler inline the calls to a specific function by using the `inline` function modifier. This keyword is best suited to small functions that are called inside loops. It should not be used on larger functions since inlining these can severely increase the size of the code, which will instead decrease performance.

```
inline int myInc(int i) { return ++i; }
```

Note that the `inline` keyword is only a recommendation. The compiler may, in its attempts to optimize the code, choose to ignore this recommendation, and it may also inline functions that do not have the `inline` modifier. Modern compilers are very good at automatically determining which functions to inline.

Auto and Decltype

Two keywords were introduced in C++11: `auto` and `decltype`. Both of these keywords are used for type deduction during compilation. The `auto` keyword works as a placeholder for a type and instructs the compiler to automatically deduce the type of the variable based on its initializer.

```
auto i = 5;      // int
auto d = 3.14;   // double
auto b = false;  // bool
```

The `auto` keyword translates to the core type of the initializer, which means that any reference and constant specifiers are dropped.

```
const int& iRef = i;
auto myAuto = iRef; // int
```

Dropped specifiers can be manually reapplied as needed. The ampersand here creates a regular (lvalue) reference.

```
const auto& myRef = iRef; // const int&
```

Alternatively, two ampersands can be used. This normally designates an rvalue reference, but in the case of `auto`, it makes the compiler automatically deduce either an rvalue or an lvalue reference, based on the given initializer.

```
int i = 1;
auto&& a = i; // int& (lvalue reference)
auto&& b = 2; // int&& (rvalue reference)
```

The `auto` specifier may be used anywhere a variable is declared and initialized. For instance, the type of the following for loop iterator is set to `auto`, since the compiler can easily deduce the type. Note that the iterator is specified as a reference. This gives better performance as it prevents copies from being made when looping over elements of a potentially large object.

CHAPTER 11 FUNCTIONS

```
#include <iostream>
#include <vector>
using namespace std;
// ...
vector<int> myVector { 1, 2, 3 };
for (auto& x : myVector) { cout << x; } // "123"
```

Prior to C++11 there was no range-based for loop or auto specifier. Iterating over a vector then required a more verbose syntax, such as the one shown here.

```
for(vector<int>::size_type i = 0; i != myVector.size(); i++) {
    cout << myVector[i]; // "123"
}
```

The `decltype` specifier works similar to `auto`, except that it deduces the exact declared type of a given expression, including references. This expression is specified in parentheses.

```
int i = 1;
int& myRef = i;
auto a = myRef; // int
decltype(myRef) b = myRef; // int&
```

In C++14, `auto` may be used as the expression for `decltype`. The keyword `auto` is then replaced with the initializing expression, allowing the exact type of the initializer to be deduced.

```
decltype(auto) c = myRef; // int&
```

Using `auto` is often the simpler choice when an initializer is available. `decltype` is mainly used to forward function return types, without having to consider whether it is a reference or value type.

```
decltype(5) getFive() { return 5; } // int
```

C++11 added a trailing return type syntax, which allows a function's return value to be specified after the parameter list, following an arrow (->). This enables the parameters to be used when deducing the return type with `decltype`. The use of `auto` in this context just means that the trailing return type syntax is being used.

```
auto getValue(int x) -> decltype(x) { return x; } // int
```

The ability to use `auto` for return type deduction was added in C++14. This enabled the core return type to be deduced directly from the return statement.

```
auto getValue(int x) { return x; } // int
```

Moreover, `auto` can be used together with `decltype` to deduce the exact type following the rules of `decltype`. This is mainly useful in the context of generic programming with templates, when there are types that are not known until runtime.

```
decltype(auto) getRef(int& x) { return x; } // int&
```

The main use for type deduction is to reduce the verbosity of the code and improve readability, particularly when declaring complicated types where the type is either difficult to know or difficult to write. Keep in mind that in modern IDEs, you can hover the mouse cursor over a variable to check its type, even if the type has been automatically deduced.

Returning Multiple Values

A convenient way to return multiple values from a function is to use a tuple. Tuples are objects that pack elements of different types into a single object.

```

#include <tuple>
#include <iostream>
using namespace std;

tuple<int, double, char> getTuple()
{
    return tuple<int, double, char>(5, 1.2, 'b');
}

```

The function can be simplified using the `auto` keyword and the `std::make_tuple` function. This function automatically deduces the types based on the provided arguments and returns a tuple.

```

auto getTuple()
{
    return make_tuple(5, 1.2, 'b');
}

```

Individual tuple elements can be extracted with the `std::get` function. Angle brackets (`<>`) are used to specify the index for the element to be retrieved. Alternatively, the type name can be used to retrieve the element if there is only one element of that type.

```

int main()
{
    auto mytuple = getTuple();
    cout << get<0>(mytuple) // "5"
          << get<char>(mytuple); // "b"
}

```

Another way to unpack a tuple is with the `std::tie` function, which will bind one or more tuple elements to the provided arguments. The `std::ignore` placeholder can be used to skip certain elements of the tuple.

```

int main()
{
    int i;
    double d;

    // Unpack tuple into variables
    tie(i, d, ignore) = getTuple();
    cout << i << " " << d; // "5 1.2"
}

```

A feature called *structured bindings* was added in C++17, providing special language support for packing and unpacking tuple-like objects. With this introduction, the `std::make_tuple` function is replaced with the following more concise code.

```

auto getTuple()
{
    return tuple(5, 1.2, 'b');
}

```

Unpacking the elements is likewise simplified and no longer requires the `std::tie` function. Note that the variables are declared automatically.

```

int main()
{
    auto [i, d, c] = getTuple();
    cout << i; // "5"
}

```

Lambda Functions

C++11 adds the ability to create lambda functions, which are unnamed function objects. This provides a compact way to define functions at their point of use, without having to create a named function or function object somewhere else. The following example creates a lambda that accepts two `int` arguments and returns their sum.

```
auto sum = [](int x, int y) -> int
{
    return x + y;
};

cout << sum(2, 3); // "5"
```

Including the return type is optional if the compiler can deduce the return value from the lambda. In C++11, this required the lambda to contain just a single return statement, whereas C++14 extended return type deduction to any lambda function. Note that the arrow (`->`) is also omitted when leaving out the return type.

```
auto sum = [](int x, int y) { return x + y; };
```

C++11 requires lambda parameters to be declared with concrete types. This requirement was relaxed in C++14, allowing lambdas to use auto type deduction. These are called generic lambda expressions.

```
auto sum = [](auto x, auto y) { return x + y; };
```

Lambdas are typically used for specifying simple functions that are only referenced once, often by passing the function object as an argument to another function. This can be done using a function wrapper with a matching parameter list and return type, as in the following example.


```

#include <iostream>
#include <functional>
using namespace std;

void call(int arg, function<void(int)> func) {
    func(arg);
}

int main() {
    auto printSquare = [](int x) { cout << x*x; };
    call(2, printSquare); // "4"
}

```

All lambdas start with a set of square brackets, called the *capture clause*. This clause specifies variables from the surrounding scope that can be used within the lambda body. This effectively passes additional arguments to the lambda, without the need to specify these in the parameter list of the function wrapper. The previous example can therefore be rewritten in the following way.

```

void call(function<void()> func) { func(); }

int main() {
    int i = 2;
    auto printSquare = [i]() { cout << i*i; };
    call(printSquare); // "4"
}

```

The variable here is captured by value, and so a copy is used within the lambda. Variables can also be captured by reference using the familiar ampersand prefix. Note that the lambda here is defined and called in the same statement.

```
int a = 1;
[&a](int x) { a += x; }(2);
cout << a; // "3"
```

It is possible to specify a default capture mode at the start of the capture clause, to indicate how any unspecified variable used inside the lambda is to be captured. A [=] means that such variables are captured by value and [&] captures them by reference. Variables captured by value are normally constant, but the mutable specifier can be used to allow such variables to be modified.

```
int a = 1, b = 1;
[&, b]() mutable { b++; a += b; }();
cout << a << b; // "31"
```

As of C++14, variables may also be initialized inside the lambda capture clause. Such variables will be type deduced as if they were declared with `auto`. Note that the parameter list following the capture clause may be omitted, as done here, provided that it is empty and the mutable specifier is not used.

```
int a = 1;
[&, b = 2] { a += b; }();
cout << a; // "3"
```

A lambda that does not capture any variables is called *stateless*. C++20 added the ability to make stateless lambdas default constructible and assignable, making the following example valid.

```
auto x = [] { return 3; };

// Default construct new lambda of same type
decltype(x) y; // valid in C++20
```

```
// Make copy of lambda
auto copy = x;

// Assign copy to x since they have same type
x = copy; // valid in C++20
```

Another feature introduced in C++20 was the ability to use lambdas in unevaluated contexts, most notably as the expression for a `decltype` specifier.

```
// Default construct inlined lambda
decltype([]{ return 3; }) a; // valid in C++20
```

CHAPTER 12

Classes

A class is a template used to create objects. To define a class, you use the keyword `class` followed by a name, a code block, and a semicolon.

A common naming convention for classes is to use mixed case, meaning that each word is initially capitalized.

```
class MyRectangle {};
```

Class members can be declared inside the class; the two main kinds are fields and methods. *Fields* are variables and they hold the state of the object. *Methods* are functions and they define what the object can do.

```
class MyRectangle
{
    int x, y;
};
```

Class Methods

A method belonging to a class is normally declared as a prototype inside of the class, and the actual implementation is placed after the class's definition. The method's name outside the class then needs to be prefixed with the class name and the scope resolution operator in order to designate to which class the method definition belongs.

```
class MyRectangle
{
    int x, y;
    int getArea();
};

int MyRectangle::getArea() { return x * y; }
```

Inline Methods

If the method is short and you want to recommend to the compiler that the method's code should be inserted (inlined) into the caller's code, one way to do this is to include the `inline` keyword in the method's definition.

```
inline int MyRectangle::getArea() { return x * y; }
```

A more convenient way is to simply define the method inside of the class. This will implicitly recommend to the compiler that the method should be inlined.

```
class MyRectangle
{
    int x, y;
    int getArea() { return x * y; }
};
```

Object Creation

The class definition is now complete. In order to use it, you first have to create an object of the class, also called an *instance*. This can be done in the same way that variables are declared.

```
int main()
{
    MyRectangle r; // object creation
}
```

Accessing Object Members

Before the members that this object contains can be accessed, they need to be declared as public in the class definition, by using the `public` keyword followed by a colon. Without this keyword, the members will have private access by default, making them inaccessible outside of the class definition.

```
class MyRectangle
{
    public:
        int x, y;
        int getArea() { return x * y; }
};
```

The members of this object can now be reached using the dot operator (`.`) after the instance name.

```
r.x = 10;
r.y = 5;
int z = r.getArea(); // 50 (5*10)
```

Any number of objects can be created based on a class, and each one of them will have its own set of fields and methods.

```
MyRectangle r2; // another instance of MyRectangle
r2.x = 25; // not same as r.x
```

When using an object pointer, the arrow operator (->) allows access to the object's members. This operator behaves like the dot operator, except that it dereferences the pointer first. It is used exclusively with pointers to objects.

```
MyRectangle r;
MyRectangle *p = &r; // object pointer
r.x = 2;
r.y = 3;
p->getArea(); // 6 (2*3)
(*p).getArea(); // alternative syntax
```

Forward Declaration

Classes, just like functions, must be declared before they can be referenced. If a class definition does not appear before the first reference to that class, a class prototype can be specified above the reference instead.

```
class MyClass; // class prototype
```

This forward declaration allows the class to be referenced in any context that does not require the class to be fully defined.

```
class MyClass; // class prototype
// ...
MyClass* p; // allowed
MyClass f(MyClass&); // allowed
MyClass o; // error, definition required
sizeof(MyClass); // error, definition required
```

Note that even with a prototype, you still cannot create an object of a class before it has been defined.

CHAPTER 13

Constructors

In addition to fields and methods, a class can contain a *constructor*. This is a special kind of method used to construct, or *instantiate*, the object. It always has the same name as the class and does not have a return type. To be accessible from another class, the constructor needs to be declared in a section marked with the public access modifier.

```
class MyRectangle
{
    public:
        int x, y;
        MyRectangle();
};
```

```
MyRectangle::MyRectangle() { x = 10; y = 5; }
```

When a new instance of this class is created, the constructor method will be called, which in this case assigns default values to the fields.

```
int main()
{
    MyRectangle s;
}
```


Constructor Overloading

As with any other function or method, the constructor can be overloaded. This will allow an object to be created with different argument lists.

```
class MyRectangle
{
public:
    int x, y;
    MyRectangle();
    MyRectangle(int, int);
};

MyRectangle::MyRectangle() { x = 10; y = 5; }
MyRectangle::MyRectangle(int a, int b) { x = a; y = b; }
```

With the two constructors defined here, the object can be initialized either with no arguments or with two arguments used to assign the fields.

```
// Calls parameterless constructor
MyRectangle r;

// Calls constructor accepting two integers
MyRectangle t(2,3);
```

C++11 added the ability for constructors to call other constructors. Using this feature, called constructor delegation, the parameterless constructor created earlier is redefined here to call the second constructor.

```
MyRectangle::MyRectangle() : MyRectangle(10, 5) {}
```

This Keyword

Inside the constructor, as well as in any other methods belonging to the object—so-called *instance methods*—a special keyword called `this` can be used. This is a pointer to the current instance of the class. It can be useful if, for example, the constructor's parameter names are the same as the field names. The fields can then still be accessed by using the `this` pointer, even though they are overshadowed by the parameters.

```
MyRectangle::MyRectangle(int x, int y)
{
    this->x = x;
    this->y = y;
}
```

Field Initialization

As an alternative to assigning fields inside the constructor, fields may also be assigned by using the *constructor initializer list*. This list starts with a colon after the constructor parameters, followed by calls to the field's own constructors.

```
MyRectangle::MyRectangle(int a, int b) : x(a), y(b) {}
```

Fields can also be assigned an initial value in their class definition, a convenient feature that was added in C++11. This is the recommended way of assigning default values to fields. The value is automatically assigned when a new instance is created, before the constructor is run. As such, this assignment can be used to specify a default value for a field that may be overridden in the constructor.

```
class MyRectangle
{
    public:
        int x = 10;
        int y = 5;
};
```

Recall that a reference must be set at the same time as it is declared. Therefore, a reference field cannot be set in the body of the constructor, but must be initialized either in the class definition or in the constructor initializer list.

```
class Foo
{
    public:
        int x;
        int& ref1 = x;
        int& ref2;
        Foo();
};
Foo::Foo() : ref2(x) {}
```

Default Constructor

If no constructors are defined for a class, the compiler will automatically create a default parameterless constructor when the program compiles. Because of this, a class can be instantiated even if no constructor has been implemented. The default constructor will only allocate memory for the object. It will not initialize the fields. In contrast to global variables, fields in C++ are not automatically initialized to their default values. The fields will contain whatever garbage is left in their memory locations until they are explicitly assigned values.

Destructor

In addition to constructors, a class can also have an explicitly defined *destructor*. This is used to release any resources allocated by the object. It is called automatically before an object is destroyed, either when the object passes out of scope or when it is explicitly deleted for objects created with the `new` operator. The name of the destructor is the same as the class name, but preceded by a tilde (`~`). A class may only have one destructor, and it never takes any arguments or returns anything.

```
class Semaphore
{
    bool *sem;
public:
    Semaphore() { sem = new bool; }
    ~Semaphore() { delete sem; }
};
```

Special Member Functions

The default constructor and the destructor are both special member functions that the compiler will automatically provide for any class that does not explicitly define them. Four other special functions are the move constructor, the move assignment operator, the copy constructor, and the copy assignment operator. With the C++11 standard came ways of controlling whether to allow these special member functions through the `delete` and `default` specifiers. The `delete` specifier forbids the calling of a function, whereas the `default` specifier explicitly states that the compiler-generated default will be used.

```

class A
{
public:
    // Explicitly include default constructor
    A() = default;

    // Explicitly include default destructor
    ~A() = default;

    // Disable move constructor
    A(A&&) noexcept = delete;

    // Disable move assignment operator
    A& operator=(A&&) noexcept = delete;

    // Disable copy constructor
    A(const A&) = delete;

    // Disable copy assignment operator
    A& operator=(const A&) = delete;
};

```

Object Initialization

C++ provides a number of different ways to create objects and initialize their fields. The following class will be used to illustrate these methods.

```

class MyClass
{
public:
    int i;
    MyClass() = default;
    MyClass(int x) : i(x) {}
};

```

Direct Initialization

The object creation syntax that has been used so far is called *direct initialization*. This syntax can include a set of parentheses that are used to pass arguments to a constructor in the class. If the parameterless constructor is used, the parentheses are left off.

```
// Direct initialization
MyClass a(5);
MyClass b;
```

Value Initialization

An object can also be *value initialized*. The object is then created by using the class name followed by a set of parentheses. The parentheses can supply constructor arguments or remain empty to construct the object using the parameterless constructor. A value initialization creates only a temporary object, which is destroyed at the end of the statement. To preserve the object, it must either be copied to another object or assigned to a reference. Assigning the temporary object to a reference will maintain the object until that reference goes out of scope.

```
// Value initialization
const MyClass& a = MyClass();
MyClass&& b = MyClass(); // alternative
```

A value-initialized object is almost identical to one created by using default initialization. A minor difference is that non-static fields will in some cases be initialized to their default values when using value initialization.

Copy Initialization

If an existing object is assigned to an object of the same type when it is declared, the new object will be *copy initialized*. This means that each member of the existing object will be copied to the new object.

```
// Copy initialization
```

```
MyClass a = MyClass(); // copy temporary object to a
MyClass b = a; // copy object a to b
```

This works because of the implicit *copy constructor* that the compiler provides, which is called for these kinds of assignments. The copy constructor takes a single argument, usually a const reference of its own type, and then constructs a copy of the specified object. Note that this behavior is different from many other languages, such as Java and C#. In those languages, initializing an object with another object will only copy the object's reference and not create a new object copy. The copy constructor can be user defined, allowing the developer to decide how the object members should be copied.

New Initialization

An object can be initialized through dynamic memory allocation by using the new keyword. Dynamically allocated memory must be used through a pointer or reference. The new operator returns a pointer, so to assign it to a reference, it needs to be dereferenced first. Keep in mind that dynamically allocated memory must be explicitly freed once it is no longer needed.

```
// New initialization
```

```
MyClass* a = new MyClass(); // object pointer
MyClass& b = *new MyClass(); // object reference
// ...
delete a;
delete &b;
```

Aggregate Initialization

There is a syntactical shortcut available when initializing an object called *aggregate initialization*. This syntax allows fields to be set by using a curly bracket–enclosed list of initializers, in the same way as can be done with arrays. Aggregate initialization can only be used when the class type does not include any constructors, virtual functions, or base classes. The fields must also be public, unless they are declared as static. Each field will be set in the order they appear in the class.

```
// Aggregate initialization
MyClass a = { 2 }; // i is 2
```

Uniform Initialization

The uniform initialization was introduced in C++11 to provide a consistent way to initialize types that work the same for any type. This syntax looks the same as aggregate initialization, without the use of the equals sign.

```
// Uniform initialization
MyClass a { 3 }; // i is 3
```

This initialization syntax works not just for classes but for any type, including primitives, strings, arrays, and standard library containers such as vector.

```
#include <string>
#include <vector>
using namespace std;

int main()
{
    int i { 1 };
    string s { "Hello" };
}
```



```

int a[] { 1, 2 };
int *p = new int [2] { 1, 2 };
vector<string> box { "one", "two" };
}

```

Uniform initialization can be used to call a constructor. This is done automatically by passing along the proper arguments for that constructor within the curly brackets.

```

// Call parameterless constructor
MyClass b {};

// Call copy constructor
MyClass c { b };

```

A class can define an initializer-list constructor. This constructor is called during uniform initialization and takes priority over other forms of construction, provided that the type specified for the `initializer_list` template matches the type of the curly bracket–enclosed list of arguments. The argument list can be any length, but all elements must be of the same type. In the following example, the type of `initializer_list` is `int`, and so the integer list used to construct this object is passed to the constructor. These integers are then displayed using a range-based for loop.

```

#include <iostream>
using namespace std;

class NewClass
{
public:
    NewClass(initializer_list<int>);
};

```

```

NewClass::NewClass(initializer_list<int> args)
{
    for (auto x : args)
        cout << x << " ";
}

int main()
{
    NewClass a { 1, 2, 3 }; // "1 2 3"
}

```

Designated Initializers

The C++20 standard introduced designated initializers, allowing any non-static field to be assigned by name in a brace-enclosed initialization list. Fields that are left unspecified will be assigned their default value, as seen in the following example.

```

class TestClass
{
public:
    int a = 1;
    int b = 2;
};

int main()
{
    TestClass o1 { .a = 3, .b = 4 }; // ok, a = 3, b = 4
    TestClass o2 { .a = 5 }; // ok, a = 5, b = 2
    TestClass o3 { .b = 6 }; // ok, a = 1, b = 6
}

```

Designated initializers can be used together with both uniform and aggregate initialization. All designated fields must appear in the order of their declaration in the class, and mixing designated and nondesignated initializers is not allowed.

```
int main()
{
    TestClass o4 { .b = 0, .a = 1 }; // error, out of order
    TestClass o5 { .a = 5, 3 }; // error, designated and non-
                                designated
}
```

CHAPTER 14

Inheritance

Inheritance allows a class to acquire the members of another class. In the following example, Square inherits from Rectangle. This is specified after the class name by using a colon followed by the `public` keyword and the name of the class to inherit from.

```
class Rectangle
{
    public:
        int x, y;
        int getArea() { return x * y; }
};

class Square : public Rectangle {};
```

Rectangle here becomes a base class of Square, which in turn becomes a derived class of Rectangle. In addition to its own members, Square gains all accessible members in Rectangle, except for its constructors and destructor.

Upcasting

An object can be upcast to its base class, because it contains everything that the base class contains. An *upcast* is performed by assigning the object to either a reference or a pointer of its base class type. In the following example, a Square object is upcast to Rectangle. When using Rectangle's

interface, the Square object will be viewed as a Rectangle, so only Rectangle's members can be accessed.

```
Square s;
Rectangle& r = s; // reference upcast
Rectangle* p = &s; // pointer upcast
```

A derived class can be used anywhere a base class is expected. For example, a Square object can be passed as an argument to a function that expects a Rectangle object. The derived object will then implicitly be upcast to its base type.

```
void setXY(Rectangle& r)
{
    r.x = 2;
    r.y = 3;
}

int main()
{
    Square s;
    setXY(s);
}
```

Downcasting

A Rectangle reference or pointer that points to a Square object can be downcast back to a Square object. This downcast has to be made explicit since downcasting an actual Rectangle to a Square is not allowed and may crash the program at runtime.

```
Square& a = static_cast<Square&>(r); // reference downcast
Square* b = static_cast<Square*>(p); // pointer downcast
```

Constructor Inheritance

To make sure the fields in the base class are properly initialized, the parameterless constructor of the base class is automatically called when an object of the derived class is created.

```
#include <iostream>
using namespace std;

class B1
{
public:
    int x;
    B1() : x(5) {}
};

class D1 : public B1 {};

int main()
{
    D1 d; // calls parameterless constructors of D1 and B1
    cout << d.x; // "5"
}
```

If there is no default constructor in the base class, the derived class must call an appropriate base class constructor. The call to the base constructor can be made explicitly from the derived constructor, by placing it in the constructor's initializer list. This allows arguments to be passed along to the base constructor.

```
class B2
{
public:
    int x;
    B2(int a) : x(a) {}
};
```

```

class D2 : public B2
{
    public:
        D2(int i) : B2(i) {} // call base constructor
};

```

An alternative solution in this case is to inherit the constructor. As of C++11, this can be done through a using statement.

```

class D2 : public B2
{
    public:
        using B2::B2; // inherit all constructors from B2
        int y { 0 };
};

int main()
{
    D2 d(3); // call inherited B2 constructor
    cout << d.x; // "3"
}

```

Note that the base class constructor cannot initialize fields defined in the derived class. Therefore, any fields declared in the derived class should initialize themselves. This is done here using the uniform notation.

Multiple Inheritance

C++ allows a derived class to inherit from more than one base class. This is called *multiple inheritance*. The base classes are then specified in a comma-separated list.

```
class Person {};  
class Employee {};  
  
class Teacher: public Person, public Employee {};
```

Multiple inheritance is not commonly used since most real-world relationships can be better described by single inheritance. It also tends to significantly increase the complexity of the code.

CHAPTER 15

Overriding

A new method in a derived class can redefine a method in a base class in order to give it a new implementation.

Hiding Derived Members

In the following example, `Rectangle`'s `getArea` method is redeclared in `Triangle` with the same signature. The signature includes the name, parameter list, and return type of the method.

```
class Rectangle
{
    public:
        int x, y;
        Rectangle(int x, int y) : x(x), y(y) {}
        double getArea() { return x * y; }
};

class Triangle : public Rectangle
{
    public:
        Triangle(int a, int b) : Rectangle(a,b) {}
        double getArea() { return x * y / 2; }
};
```

If a `Triangle` object is created and the `getArea` method is invoked, then `Triangle`'s version of the method will get called.

```
Triangle t { 2,3 }; // uniform initialization
t.getArea(); // 3 (2*3/2) calls Triangle's version
```

However, if the `Triangle` is upcast to a `Rectangle`, then `Rectangle`'s version will get called instead.

```
Rectangle& r = t; // upcast
r.getArea(); // 6 (2*3) calls Rectangle's version
```

That is because the redefined method has only hidden the inherited method. This means that `Triangle`'s implementation is redefined downward in the class hierarchy to any child classes of `Triangle`, but not upward to the base class.

Overriding Derived Members

In order to redefine a method upward in the class hierarchy—what is called *overriding*—the method needs to be declared with the `virtual` modifier in the base class. This modifier allows the method to be overridden in derived classes.

```
class Rectangle
{
    public:
        int x, y;
        virtual int getArea() { return x * y; }
};
```

Calling the `getArea` method from `Rectangle`'s interface will now invoke `Triangle`'s implementation. This is called polymorphism—when a method call causes a different method to be executed depending on the

type of object that invokes the method. Note that polymorphism requires the use of references or pointers.

```
Triangle t { 2,3 };
Rectangle& r = t;
r.getArea(); // 3 (2*3/2) calls Triangle's version
```

C++11 added the `override` specifier, which indicates that a method is intended to replace an inherited method. Using this specifier allows the compiler to check that there is a virtual method with that same signature. This prevents the possibility of accidentally creating a new virtual method in a derived class. It is recommended to always include this specifier when overriding methods.

```
class Triangle : public Rectangle
{
public:
    virtual double getArea() override { return x * y / 2; }
};
```

Another specifier introduced in C++11 is `final`. This specifier prevents a virtual method from being overridden in derived classes. It also prevents derived classes from using that same method signature.

```
class Base
{
    virtual void foo() final {}
};

class Derived : public Base
{
    void foo() {} // error: Base::foo marked as final
};
```

The `final` specifier can also be applied to a class to prevent any class from inheriting it.

```
class B final {};  
class D : B {}; // error: B marked as final
```

Base Class Scoping

It is still possible to access a redefined method from a derived class by typing the class name followed by the scope resolution operator. This is called *base class scoping* and can be used to allow access to redefined methods that are any number of levels deep in the class hierarchy.

```
class Triangle : public Rectangle  
{  
    public:  
        Triangle(int a, int b) { x = a; y = b; }  
        int getArea() override { return Rectangle::getArea() / 2; }  
};
```

Pure Virtual Functions

Sometimes a base class knows that all derived classes must implement a certain method, but the base class cannot provide a default implementation for that method. The base class can then declare the method as a pure virtual function, by assigning it the value zero, in order to force deriving classes to implement this method.

```
class Shape  
{  
    public:  
        virtual double getArea() = 0; // pure virtual function  
};
```

A class with one or more pure virtual functions is called an abstract class since it is incomplete and therefore cannot be instantiated. Abstract classes are mainly used for upcasting, so that deriving classes can use its interface through a pointer or reference type.

```
#include <iostream>
class Rectangle : public Shape
{
    public:
        int x = 1, y = 2;
        virtual int getArea() override { return x * y; }
};

void printArea(Shape& s) {
    std::cout << s.getArea();
}

int main()
{
    Rectangle r;
    printArea(); // "2"
}
```

A class consisting of only pure virtual functions is known as an interface. Such a class is functionally the same as an interface in other languages such as C# or Java.

CHAPTER 16

Access Levels

Every class member has an accessibility level that determines where the member will be visible. There are three of them available in C++: `public`, `protected`, and `private`. The default access level for class members is `private`. To change the access level for a section of a class, an access modifier is used, followed by a colon. Every field or method that comes after this label will have the specified access level, until another access level is set or the class declaration ends.

```
class MyClass
{
    int myPrivate;

public:
    int myPublic;
    void publicMethod();
};
```

Private Access

All members regardless of their access level are accessible in the class in which they are declared, which is called the enclosing class. This is the only place where private members can be accessed.

```

class MyClass
{
    // Unrestricted access
    public: int myPublic;

    // Defining or derived class only
    protected: int myProtected;

    // Defining class only
    private: int myPrivate;

    void test()
    {
        myPublic = 0; // allowed
        myProtected = 0; // allowed
        myPrivate = 0; // allowed
    }
};

```

Protected Access

A protected member can also be accessed from inside a derived class, but it cannot be reached from an unrelated class.

```

class MyChild : public MyClass
{
    void test()
    {
        myPublic = 0; // allowed
        myProtected = 0; // allowed
        myPrivate = 0; // inaccessible
    }
};

```

Public Access

Public access gives unrestricted access from anywhere in the code.

```
class OtherClass
{
    void test(MyClass& c)
    {
        c.myPublic = 0; // allowed
        c.myProtected = 0; // inaccessible
        c.myPrivate = 0; // inaccessible
    }
};
```

Access Level Guideline

As a guideline, when choosing an access level, it is generally best to use the most restrictive level possible. This is because the more places a member can be accessed, the more places it can be accessed incorrectly, which makes the code harder to debug. Using restrictive access levels will also make it easier to modify the class without breaking the code for any other programmers using that class.

When coding in the real world, fields should always be private and only exposed through public or protected getter and setter methods. This makes it easier to ensure that fields are accessed correctly, as the setter can check that the assigned value is valid for the specific field. By leaving out either the getter or setter method, a field may also be restricted to only write or read access from outside the class.


```
class Person
{
    private:
        int age;

    public:
        // Setter
        void setAge(int a)
        {
            if (age > 200) age = 200;
            else if (age < 0) age = 0;
            else age = a;
        }
        // Getter
        int getAge()
        {
            return age;
        }
};
```

Friend Classes and Functions

A class can be allowed to access the private and protected members of another class by declaring that class a friend. This is done by using the friend modifier. The friend is allowed to access all members in the class where the friend is defined, but not the other way around.

```
class MyClass
{
    int myPrivate;
```

```

    // Give OtherClass access
    friend class OtherClass;
};

class OtherClass
{
    void test(MyClass& c) {
        c.myPrivate = 0; // allowed
    }
};

```

Likewise, a method of another class may be marked as a friend to allow it to access all members in the defining class.

```

class MyClass;
class OtherClass
{
    public:
        void test(MyClass& c);
        void test2(MyClass& c);
};

class MyClass
{
    int myPrivate;
    friend void OtherClass::test(MyClass&);
};

void OtherClass::test(MyClass& c) {
    c.myPrivate = 0; // allowed
}

void OtherClass::test2(MyClass& c) {
    c.myPrivate = 0; // not allowed
}

```

A global function can also be declared as a friend to a class in order to gain the same level of access.

```
class MyClass
{
    int myPrivate;

    // Give myFriend access
    friend void myFriend(MyClass& c);
};

void myFriend(MyClass& c) {
    c.myPrivate = 0; // allowed
}
```

Public, Protected, and Private Inheritance

When a class is inherited in C++, it is possible to change the access level of the inherited members. Public inheritance allows all members to keep their original access level. Protected inheritance reduces the access of public members to protected. Private inheritance restricts all inherited members to private access.

```
class MyChild : private MyClass
{
    // myPublic is private
    // myProtected is private
    // myPrivate is private
};
```

Private is the default inheritance level, although public inheritance is the one that is nearly always used.

CHAPTER 17

Static

The `static` keyword is used to create class members that exist in only one copy, which belongs to the class itself. These members are shared among all instances of the class. This is different from instance (non-static) members, which are created as new copies for each new object.

Static Fields

A static field (class field) is initialized outside of the class declaration. This initialization will take place only once, and the static field will then remain initialized throughout the life of the application.

```
class MyCircle
{
    public:
        double r; // instance field (one per object)
        static double pi; // static field (only one copy)
};

double MyCircle::pi = 3.14159;
```

To access a static member from outside the class, the name of the class is used followed by the scope resolution operator and the static member's name. This means that there is no need to create an instance of a class in order to access its static members.

```
int main()
{
    double p = MyCircle::pi;
}
```

There are two exceptions to the rule that all fields must be initialized outside of the class. First exception is if the static field is of an integral or enum type and it is declared as a constant, using the `const` modifier. Second exception is if the field uses the `inline` modifier, a feature which was introduced in C++17.

```
class MyClass
{
    static inline double myDouble = 1.23;
    static const int myInt = 1;
};
```

Static Methods

In addition to fields, methods can also be declared as `static`, in which case they can also be called without having to create an instance of the class. However, because a static method is not part of any instance, it cannot use instance members as it does not have an implicit `this` pointer. Methods should therefore only be declared `static` if they perform a generic function that is independent of any instance variables. Instance methods, in contrast to static methods, can use both static and instance members.

```
class MyCircle
{
public:
    double r;
    static inline double pi = 3.14159;
```

```

double getArea() { return pi * r * r; }
static double newArea(double a) { return pi * a * a; }
};

int main()
{
    double a = MyCircle::newArea(1);
}

```

Static Local Variables

Local variables inside a function can be declared as `static` to make the function remember the variable for the lifetime of the application. A static local variable is only initialized once when execution first reaches the declaration, and that declaration is then ignored every subsequent time the execution passes through.

```

void myFunc()
{
    static int count = 0; // holds # of calls to function
    count++;
}

```

Static Global Variables

One last place where the `static` keyword can be applied is to global variables. This will limit the accessibility of the variable to only the current source file and can therefore be used to help avoid naming conflicts.

```

// Only visible within this source file
static int myGlobal;

```

This application of static is seldom used. The preferred way to limit code entities to a single source file is to enclose them in an unnamed namespace.

```
namespace
{
    // Only visible within this source file
    int myGlobal;
}
```

CHAPTER 18

Enum Types

An enum is a user-defined type consisting of a fixed list of named constants. In the following example, the enumeration type is called `Color`, and it contains three constants: `red`, `green`, and `blue`.

```
enum class Color { red, green, blue };
```

The `Color` type can be used to create variables that may hold one of these constant values. Enum class constants must be prefixed with the enum name as seen here.

```
int main()
{
    Color c = Color::red;
}
```

Enum Example

The `switch` statement provides a good example of when enumerations can be useful. Compared to using ordinary constants, the enumeration has the advantage in that it allows the programmer to clearly specify what values a variable should be allowed to contain.


```
switch(c)
{
    case Color::red:    break;
    case Color::green: break;
    case Color::blue:  break;
}
```

Enum Constant Values

Usually, there is no need to know the underlying values that the constants represent, but in some cases, it can be useful. By default, the first constant in the enum list has the value zero, and each successive constant is one value higher.

```
enum class Color
{
    red,    // 0
    green,  // 1
    blue    // 2
};
```

These default values can be overridden by assigning values to the constants. The values can be computed and do not have to be unique. A constant that is not assigned a value will have a value one higher than the previous assigned enum value.

```
enum class Color
{
    red = 5, // 5
    green = red, // 5
    blue = green + 2, // 7
    yellow // 8
};
```

Enum Scope

An enum does not have to be declared globally. It can also be placed within a class as a class member or locally within a function.

```
class MyClass
{
    enum class Color { red, green, blue };
};

void myFunction()
{
    enum class Color { red, green, blue };
}
```

Weakly Typed Enums

The enum class type described so far was introduced in C++11 to provide a safer alternative to the weakly typed enum inherited from C. This legacy enum is defined in the same way as the enum class, but without the `class` keyword.

```
// Weakly typed enum
enum Speed
{
    fast,
    normal,
    slow
};
```

With this weakly typed enum, the specified constants do not belong within the scope of the enum name. Such an enum constant can therefore be referenced even without qualifying it with the enum name.

```
Speed s1 = fast;
Speed s2 = Speed::normal;
```

It is preferable to use enum classes rather than weakly typed enums because of their type safety and because their constants are scoped to the enum name. Since enum classes are strongly typed, they will not implicitly convert to integer types.

```
// Weakly typed enum
enum Speed { fast, normal, slow };
Speed s = fast;
if (s == fast) {} // ok
if (s == 0) {} // ok

// Strongly typed enum
enum class Color { red, green, blue };
Color c = Color::red;
if (c == Color::red) {} // ok
if (c == 0) {} // error
```

C++20 added the ability to import an enum class into the local scope with a using statement. This avoids needless repetition of the enum class name, by making the enum class members accessible like regular enum members within a specific scope. Be sure not to import enums into a too large scope, else the main advantage of using strongly typed enums is lost.

```
#include <iostream>
using namespace std;
enum class Color { red, green, blue };
void colorPrint(Color c)
{
    // Import enum members to local scope
    using enum Color;
    switch (c)
```

```
{  
    case red: cout << "red";  
    case green: cout << "green";  
    case blue: cout << "blue";  
}  
}
```

Enum Constant Type

The underlying integer type of the regular enum is not defined by the standard and may vary between implementations. In contrast, an enum class always uses the int type by default. For both types of enums, the type can be overridden to another integer type, as in the following example.

```
// Enum with constant type set to unsigned short  
enum class MyEnum : unsigned short {};
```

CHAPTER 19

Structs and Unions

Structs

A struct in C++ is equivalent to a class, except that members of a struct default to public access, instead of private access as in classes. By convention, structs are used instead of classes to represent simple data structures that mainly contain public fields.

```
struct Point
{
    int x, y; // public
};

class Point
{
    int x, y; // private
};
```

Struct Initialization

To declare objects of a struct, you use the normal declaration syntax.

```
Point p, q; // object declarations
```

Another alternative syntax sometimes used with structs (and classes) is to declare the objects when the struct is defined by placing the object names before the final semicolon. This position is known as the *declarator list* and can contain a comma-separated sequence of declarators.

```
struct Point
{
    int x, y;
} r, s; // object declarations
```

When using object declarations, the name of the struct may optionally be omitted. This is called an anonymous struct.

```
struct
{
    int x, y;
} r, s;
```

Aggregate initialization is also commonly used with structs, since this syntactical shortcut only works for the aggregate types: array, class, struct, and union. For this initialization to work, the type must not include any private or protected non-static fields.

```
int main()
{
    // Aggregate initialization
    Point p = { 2, 3 };
}
```

For compilers supporting C++11 or later versions, the uniform initialization syntax is preferred, as it removes the distinction between initialization of aggregate and non-aggregate types.

```
int main()
{
    // Uniform initialization
    Point q { 2, 3 };
}
```

Union

Although similar to struct, the union type is different in that all fields share the same memory position. Therefore, the size of a union is the size of the largest field it contains. For example, in the following case, this is the integer field, which is four bytes large.

```
union Mix
{
    char c; // 1 byte
    short s; // 2 bytes
    int i; // 4 bytes
} m;
```

This means that the union type can be used to store only one value at a time, because changing one field will overwrite the values of the others.

```
int main()
{
    m.c = 0xFF; // set first 8 bits
    m.s = 0; // reset first 16 bits
}
```

The benefit of a union, in addition to efficient memory usage, is that it provides multiple ways of viewing the same memory location. For example, the following union has three data members that allow access to the same group of four bytes in multiple ways.

```

union Mix
{
    char c[4];           // 4 bytes
    struct { short hi, lo; } s; // 4 bytes
    int i;               // 4 bytes
} m;

```

The integer field will access all four bytes at once. With the struct, two bytes can be viewed at a time, and by using the char array, each byte can be referenced individually.

```

int main()
{
    // Set i = 11111111 00000000 11110000 00001111
    m.i=0xFF00F00F;
    m.s.lo; // 11111111 00000000
    m.s.hi; //                11110000 00001111
    m.c[3]; // 11111111
    m.c[2]; //                00000000
    m.c[1]; //                11110000
    m.c[0]; //                00001111
}

```

Anonymous Union

A union type can be declared without a name for the type or the object. This is called an anonymous union and defines an unnamed object whose members can be accessed directly from the scope where the object is declared. Unlike regular unions, an anonymous union cannot contain methods or nonpublic members.


```
int main()
{
    union { short s; }; // defines an unnamed union object
    s = 15;
}
```

An anonymous union that is declared globally must be made static.

```
static union {};
```

CHAPTER 20

Operator Overloading

Operator overloading allows operators to be redefined and used where one or both of the operands are of a user-defined class. When it's done correctly, this can simplify the code and make user-defined types as easy to use as the primitive types.

In the following example, there is a class called `MyNum` with an integer field and a constructor for setting that field. The class also has a method that adds two `MyNum` objects and returns the result as a new object.

```
class MyNum
{
    int val;
public:
    MyNum(int i) : val(i) {}

    MyNum add(const MyNum &a) const {
        return MyNum( val + a.val );
    }
};
```

As seen here, two `MyNum` instances can be added together using this method.

```
MyNum a = MyNum(10),
      b = MyNum(5);
MyNum c = a.add(b);
```

Binary Operator Overloading

What operator overloading does is simplify this syntax and thereby provide a more intuitive interface for the class. To convert the `add` method to an overload for the addition sign, replace the name of the method with the operator keyword followed by the operator that is to be overloaded. The whitespace between the keyword and the operator can optionally be left out.

```
MyNum operator + (const MyNum &a) const {
    return MyNum( val + a.val );
}
```

Since the class now overloads the addition sign, this operator can be used to perform the needed calculation.

```
MyNum c = a + b;
```

Keep in mind that the operator is simply an alternative syntax for calling the actual method.

```
MyNum d = a.operator+(b);
```

Unary Operator Overloading

Addition is a binary operator, because it takes two operands. The first operand is the object from which the method is called, and the second operand is that which is passed to the method. When overloading a unary operator, such as prefix increment (`++`), there is no need for a method parameter since these operators only affect the object from which they are called. With unary operators, a reference of the same type as the object should always be returned. This is because, when using a unary operator on an object, programmers expect the result to return the same object and not just a copy. On the other hand, when using a binary operator,

programmers expect a copy of the result to be returned, and therefore return by value should be used.

```
// Increment prefix
MyNum& operator++()
{
    ++val;
    return *this;
}
```

Not all unary operators should return by reference. The two postfix operators—post-increment and post-decrement—should instead return by value, because the postfix operations are expected to return the state of the object before the increment or decrement occurs. Note that the postfix operators have an unused `int` parameter specified. This parameter is used to distinguish them from the prefix operators.

```
// Increment postfix
MyNum operator++(int)
{
    MyNum t = MyNum(val);
    ++val;
    return t;
}
```

Comparison Operator Overloading

The three-way comparison operator (`<=>`) was added in C++20 to provide a simple way to overload the four comparison operators `<`, `>`, `<=`, and `>=`. When used on a whole number type, as in the following example, the operator returns an object representing either equal, less, or greater.

```

#include <compare> // std::strong_ordering
class Length
{
public:
    int i;
    std::strong_ordering operator<=>(const Length& right) const {
        return i <=> right.i;
    }
};

```

With this operator defined, the compiler automatically generates all four comparison operators based on this method.

```

int main()
{
    Length n1 { 1 }, n2 { 2 };
    bool b = n1 < n2; // true
}

```

The following example illustrates how the resulting object from the three-way comparison operator can be used.

```

#include <compare>
#include <iostream>

int main()
{
    int x = 5;
    auto result = x <=> 0;
    if (result > 0) { // true
        cout << "5 > 0";
    }
}

```

Another feature of C++20 is that the compiler will generate the inequality operator (`!=`) if the equality operator (`==`) is defined. As to be expected, the inequality operator returns the inverse of the equality operator.

```
class Length
{
public:
    int i;
    bool operator==(const Length& other) const {
        return i == other.i;
    }
};

int main()
{
    Length m1 { 1 }, m2 { 2 };
    bool b1 = m1 == m2; // false
    bool b2 = m1 != m2; // true
}
```

Any of the four comparison operators (`<`, `>`, `<=`, and `>=`), as well as equal to (`==`) and the three-way comparison operator (`<=>`), can be explicitly defaulted. This will make the compiler automatically implement the specified comparison method, which will compare the fields of the class in the order in which they are defined, stopping early when a non-equal result is found. The return type, the type of ordering, is automatically deduced based on the return type of the three-way comparison operator. If the operator is defaulted, as seen in the following, the compiler will generate all six of the comparison operators (`<`, `>`, `<=`, `>=`, `==`, and `!=`).

```
#include <compare>

class Point
{
    int x, y;
public:
    auto operator<=>(const Point&) const = default;
};

int main()
{
    Point p1 { 1, 10 }, p2 { 2, 0 };
    bool b = p1 < p2; // true (p1.x < p2.x)
}
```

Overloadable Operators

C++ allows overloading almost all operators in the language. As can be seen in the following table, most operators are of the binary type. Only a few of them are unary, and some special operators cannot be categorized as either. There are also some operators that cannot be overloaded at all.

Binary Operators	Unary Operators
+ - * / %	+ - ! ~ & * ++ --
= += -= *= /= %=	Special operators
&= ^= = <<= >>=	() [] delete new
== != > < >= <= <=>	Not overloadable
& ^ << >> &&	. * :: ?: # ## sizeof typeid alignof
-> ->* ,	noexcept

CHAPTER 21

Custom Conversions

Custom type conversions can be defined to allow an object to be constructed from or converted to another type. In the following example, there is a class called `MyNum` with a single integer field. With conversion constructors, it is possible to allow integer types to be implicitly converted to this object's type.

```
class MyNum
{
    int value;
};
```

Implicit Conversion Constructor

For this type of conversion to work, a constructor needs to be added that takes a single parameter of the desired type, in this case an `int`.

```
class MyNum
{
public:
    MyNum(int i) : value(i) {}
private:
    int value;
};
```


When an integer is assigned to an object of `MyNum` type, this constructor will implicitly be called to perform the type conversion.

```
MyNum a = 5; // implicit conversion
```

This means that any constructor that takes exactly one argument can be used both for constructing objects and for performing implicit type conversions to that object type.

```
MyNum b(5); // object construction
MyNum c = 5; // implicit conversion
```

These conversions will work not only for the specific parameter type but also for any type that can be implicitly converted to it. For example, a `char` can be implicitly converted to an `int` and can therefore be implicitly changed into a `MyNum` object as well.

```
MyNum d = 'H'; // implicit conversion (char->int->MyNum)
```

When using braced initializers, even constructors with multiple parameters can be converting constructors. In the following example, an integer list is implicitly converted to a `Point`.

```
class Point
{
public:
    Point(int x, int y) : x(x), y(y) {}
private:
    int x, y;
};

int main()
{
    Point p = { 1,2 };
}
```

Explicit Conversion Constructor

To help prevent potentially unintended object type conversions, it is possible to disable the implicit use of converting constructors. The `explicit` constructor modifier is then applied, which specifies that the constructor may only be used for object construction and not for type conversion.

```
class MyNum
{
public:
    int value;
    explicit MyNum(int i) { value = i; }
};
```

The `explicit` constructor syntax or an explicit conversion must be used when creating an object of this type.

```
MyNum a = 5; // error
MyNum b(5); // allowed
MyNum c = MyNum(5); // allowed
MyNum d = static_cast<MyNum>(5); // allowed
```

Conversion Operators

Custom conversion operators allow conversions to be specified in the other direction: from the object's type to another type. The `operator` keyword is then used, followed by the target type, a set of parentheses, and a method body. The body returns a value of the target type, in this case `int`.

```
class MyNum
{
public:
    int value;
    operator int() { return value; }
};
```

When objects of this class are evaluated in an `int` context, this conversion operator will be called to perform the type conversion.

```
MyNum a { 5 };
int i = a; // 5
```

Explicit Conversion Operators

The C++11 standard added explicit conversion operators to the language. Similar to explicit constructors, the inclusion of the `explicit` keyword prevents the conversion operator from being implicitly called.

```
class True
{
public:
    explicit operator bool() const {
        return true;
    }
};
```

This class provides a Boolean value that prevents its objects from mistakenly being used in a mathematical context through the `bool` conversion operator. In the next example, the first comparison results in a compilation error since the `bool` conversion operator cannot be implicitly called. The second comparison is allowed because the conversion operator is explicitly called through the type cast.

```
True a, b;  
if (a == b) {} // error  
if (static_cast<bool>(a) == static_cast<bool>(b)) {} // allowed
```

Bear in mind that contexts requiring a bool value, such as the condition for an if statement, count as explicit conversions.

```
if (a) {} // allowed
```

CHAPTER 22

Namespaces

Namespaces are used to avoid naming conflicts by allowing entities, such as classes and functions, to be grouped under a separate scope. In the following code, there are two classes that belong to the global scope. Since both classes share the same name and scope, the code will not compile.

```
class Table {};  
class Table {}; // error: class type redefinition
```

One way to solve this problem is to rename one of the conflicting classes. Another solution is to group one or both of them under a different namespace by enclosing each in a namespace block. The classes then belong to different scopes and so will no longer cause a naming conflict.

```
namespace furniture  
{  
    class Table {};  
}  
  
namespace html  
{  
    class Table {};  
}
```

Accessing Namespace Members

To access a member of a namespace from outside that namespace, you must specify the member's fully qualified name. This means that the member name has to be prefixed with the namespace it belongs to, followed by the scope resolution operator.

```
int main()
{
    furniture::Table fTable;
    html::Table hTable;
}
```

Nesting Namespaces

It is possible to nest namespaces any number of levels deep to further structure code entities.

```
namespace furniture
{
    namespace wood { class Table {}; }
}
```

As of C++17, the nesting of namespaces can be shortened in the following manner.

```
namespace furniture::wood { class Table {}; }
```

Ensure that the nested namespace members are qualified with the full namespace hierarchy when using them from another namespace.

```
furniture::wood::Table fTable;
```

Importing Namespaces

To avoid having to specify the namespace every time one of its members is used, the namespace can be imported into the global or local scope with the help of a `using` directive. This directive includes the `using namespace` keywords followed by the namespace to be imported. It can be placed either locally or globally. Locally, the directive will only be in scope until the end of the code block, while at the global scope, it will apply to the whole source file following its declaration.

```
using namespace html; // global namespace import

int main()
{
    using namespace html; // local namespace import
}
```

Keep in mind that importing a namespace into the global scope defeats the main purpose of using namespaces, which is to avoid naming conflicts. Such conflicts, however, are mainly an issue in projects that use several independently developed code libraries.

Namespace Member Import

If you want to avoid both typing the fully qualified name and importing the whole namespace, there is a third alternative available. That is to only import the specific members that are needed from the namespace. This is done by declaring one member at a time in a `using` declaration, which consists of the `using` keyword followed by the fully qualified namespace member to be imported.

```
using html::Table; // import a single namespace member
```

Namespace Alias

Another way to shorten the fully qualified name is to create a namespace alias. The namespace keyword is then used followed by an alias name, to which the fully qualified namespace is assigned.

```
namespace myAlias = furniture::wood; // namespace alias
```

This alias can then be used instead of the namespace qualifier that it represents.

```
myAlias::Table fTable;
```

Note that namespace aliases, as well as using directives and using declarations, may be declared either globally or locally.

Type Alias

Aliases can also be created for types. A type alias is defined with a using statement. With this syntax, the keyword using is followed by the alias name and then assigned the type.

```
using MyType = furniture::wood::Table;
```

The alias can then be used as a synonym for the specified type.

```
MyType t;
```

Before using statements were introduced in C++11, type aliases were defined with typedef. In such a statement, the typedef keyword is followed by the type name and then the alias name. Both methods for declaring aliases are equivalent, but the using statement is preferred as it is considered easier to read than the typedef statement.

```
typedef furniture::wood::Table MyType;
```


Aliases should be used with care since they may obfuscate the code. However, if used properly, a type alias can simplify a long or confusing type name. Another function they provide is the ability to change the definition of a type from a single location.

Including Namespace Members

Keep in mind that in C++ merely importing a namespace does not provide access to the members included in that namespace. In order to access the namespace members, the prototypes also have to be made available, for example, by using the appropriate `#include` directives.

```
// Include input/output prototypes
#include <iostream>

// Import standard library namespace to global scope
using namespace std;
```

CHAPTER 23

Constants

A *constant* is a variable that has a value that cannot be changed once the constant has been assigned. This allows the compiler to enforce that the variable's value is not changed anywhere in the code by mistake.

Constant Variables

A variable can be made into a constant by adding the `const` keyword either before or after the data type. This modifier means that the variable becomes read-only, and it must therefore be assigned a value at the same time as it is declared. Attempting to change the value anywhere else results in a compile-time error.

```
const int var = 5;  
int const var2 = 10; // alternative order
```

Constant Pointers

When it comes to pointers, `const` can be used in two ways. First, the pointer can be made constant, which means that it cannot be changed to point to another location.

```
int myPointee;  
int* const p = &myPointee; // pointer constant
```

Second, the pointee can be declared constant. This means that the variable pointed to cannot be modified through this pointer.

```
const int* q = &myPointee; // pointee constant
```

It is possible to declare both the pointer and the pointee as constant to make them both read-only. Reading the type from right to left makes it easier to understand, so in this case: `r` is a `const` pointer to a `const int`.

```
const int* const r = &myPointee; // pointer & pointee constant
```

Note that constant variables may not be pointed to by a non-constant pointer. This prevents programmers from accidentally rewriting a constant variable using a pointer.

```
const int myConst = 3;  
int* s = &myConst; // error: const to non-const assignment
```

Constant References

References can be declared constant in the same way as pointers. However, since reseating a reference is never allowed, declaring the reference as `const` would be redundant. It only makes sense to protect the referee from changing.

```
const int& y = myPointee; // referee constant
```

Constant Objects

Just as with variables, pointers, and references, objects can also be declared constant. Take the following class as an example.

```
class MyClass
{
    public:
        int x;
        void setX(int a) { x = a; }
};
```

A constant object of this class cannot be reassigned to another instance. The `const`-ness of an object also affects its fields and prevents them from being changed.

```
const MyClass a, b;
a = b;      // error: object is const
a.x = 10;   // error: object field is const
```

Constant Methods

Because of this last restriction, a non-constant method cannot be called on a constant object since such methods are allowed to change the object's fields.

```
a.setX(2); // error: cannot call non-const method
```

They may only call constant methods, which are methods that are marked with the `const` modifier before the method body.

```
int getX() const { return x; } // constant method
```

This `const` modifier means that the method is not allowed to modify the state of the object and can therefore safely be called from a constant object of the class. More specifically, the `const` modifier applies to the `this` pointer that is implicitly passed to the method. This effectively restricts the method from modifying the object's fields or calling any non-constant methods in the class.

Constant Return Type and Parameters

In addition to making a method constant, the return type and method parameters may also be made read-only. For example, if a field is returned by reference instead of by value from a constant method, it is important that it is returned as a constant in order to maintain the const-ness of the object. Not all C++ compilers will be able to catch this subtle mistake.

```
const int& getX() const { return x; }
```

Objects should always be passed to and returned from functions and methods by const reference. This improves performance as it prevents unnecessary copies from being made.

Constant Fields

Both static and instance fields in a class can be declared constant. A constant instance field must be assigned its value using either in-class initializers or the constructor initialization list.

```
class MyClass
{
public:
    int a;
    const int b;
    const int c = 3;
    MyClass() : a(1), b(2) {}
};
```

A constant static field has to be defined outside of the class declaration, in the same way as non-constant static fields. The exception to this is when the constant static field is either inline or of an integer data type. Such a field may also be initialized within the class at the same time as the field is declared.

```

class MyClass
{
public:
    const static double c1;
    const inline static double c2 = 1.23;
    const static int c3 = 5;
};
const double MyClass::c1 = 1.23;

```

Constant Expressions

The keyword `constexpr` was introduced in C++11 to indicate a constant expression. Like `const` it can be applied to variables to make them constant, causing a compilation error if any code attempts to modify the value.

```

constexpr int myConst = 5;
myConst = 3; // error: variable is const

```

Unlike `const` variables, which may be assigned at runtime, a constant expression variable will always be computed at compile time. Such a variable can therefore be used whenever a compile-time constant is needed, such as in array or enum declarations. Prior to C++11, this was only allowed for constant integer and enumeration types.

```

int myArray[myConst + 1]; // allowed

```

Functions and class constructors may also be defined as constant expressions, which is not allowed with `const`. Using `constexpr` on a function limits what the function is allowed to do. In short, the function can only reference other `constexpr` functions and global `constexpr` variables.

```
constexpr int getDefaultSize(int multiplier)
{
    return 3 * multiplier;
}
```

The return value for a `constexpr` function is guaranteed to be evaluated at compile time only when its arguments are constant expressions, and the return value is used where a compile-time constant is necessary.

```
// Compile-time evaluation
int myArray[getDefaultSize(10)];
```

If the function is called without constant arguments, it returns a value at runtime just like a regular function.

```
// Runtime evaluation
int mul = 10;
int size = getDefaultSize(mul);
```

As of C++17, a lambda expression is implicitly `constexpr` if it satisfies the conditions of a `constexpr` function. Such a lambda may therefore also be used in a compile-time context.

```
auto answer = [](int i) { return 10+i; };
constexpr int reply = answer(32); // "42"
```

Constructors can be declared with `constexpr`, to construct a constant expression object. Such a constructor must be trivial.

```
class Circle
{
public:
    int r;
    constexpr Circle(int x) : r(x) {}
};
```

When called with a constant expression argument, the result will be a compile-time generated object with read-only fields. With any other arguments, it will behave as an ordinary constructor.

```
// Compile-time object
constexpr Circle c1(5);

// Runtime object
int x = 5;
Circle c2(x);
```

One additional use for `constexpr` was added in C++17: the ability to evaluate conditional statements at compile time. This feature allows branches of an `if` statement to be discarded at compile time based on a constant condition, potentially reducing compilation time as well as the size of the compiled file.

```
constexpr int debug = 0;
if constexpr(debug) {
    // Discarded if condition is false
}
```

Up until C++17, virtual functions could not be defined as `constexpr`. This restriction was lifted in C++20, allowing such virtual functions to be called within a constant expression. Note that a `constexpr` virtual function can override a non-`constexpr` virtual function, as seen in the following example.

```
struct Parent {
    virtual int num() const = 0;
};

struct Child: public Parent {
    constexpr virtual int num() const { return 3; }
};
```



```
constexpr Child c;
static_assert( c.num() == 3, "num is not 3" );
```

The `static_assert` declaration seen here is used to make an assertion at compile time. If the assertion fails, which occurs when the condition evaluates to false, the compiler halts compilation and displays the error message.

Immediate Functions

As mentioned previously, the return value of a `constexpr` function is not always required to be evaluated at compile time. For such a purpose, C++20 introduced immediate functions. An immediate function is defined using the `constexpr` keyword, which designates that the function must always return a compile-time constant. Such a function can be used in a context requiring a constant expression, as seen in the following example.

```
constexpr int doubleIt(int i) {
    return 2*i;
}
constexpr int a = doubleIt(10); // ok
int x = 10;
int b = doubleIt(x); // error: call does not produce a constant
```

Constant Guideline

In general, it is a good idea to always declare variables as constants if they do not need to be modified. This ensures that the variables are not changed anywhere in the program by mistake, which in turn will help prevent bugs. There is also a performance gain by allowing the compiler the opportunity to hard-code constant expressions into the compiled program. This allows the expression to be evaluated only once—during compilation—rather than every time the program runs.

CHAPTER 24

Preprocessor

The preprocessor is a text substitution tool that modifies the source code before the compilation takes place. This modification is done according to the preprocessor directives that are included in the source files. The directives are easily distinguished from other programming code in that they start with a hash sign (#). They must always appear as the first non-whitespace character on a line, and they do not end with a semicolon. The following table shows the preprocessor directives available in C++ along with their functions.

Directive	Description
<code>#include</code>	File include
<code>#define</code>	Macro definition
<code>#undef</code>	Macro undefined
<code>#ifdef</code>	If macro defined
<code>#ifndef</code>	If macro not defined
<code>#if</code>	If
<code>#elif</code>	Else if
<code>#else</code>	Else
<code>#endif</code>	End if
<code>#line</code>	Set line number
<code>#error</code>	Abort compilation
<code>#pragma</code>	Set compiler option

Including Source Files

The `#include` directive inserts the contents of a file into the current source file. Its most common use is to include header files, both user-defined and library ones. Library header files are enclosed between angle brackets (`<>`). This tells the preprocessor to search for the header in the default directory where it is configured to look for standard header files.

```
#include <iostream> // search library directory
```

Header files that you create for your own program are enclosed within double quotes ("`"`"). The preprocessor will then search for the file in the same directory as the current file. If the header is not found there, the preprocessor will then search in the directories where it has been configured to look for header files, and after that it will look in the default folder for standard header files.

```
#include "MyFile.h" // search current, configured and default
directories
```

The double-quoted form can also be used to specify an absolute or relative path to the file, although specifying paths like this is discouraged.

```
#include "C:\MyFile.h" // absolute path
#include "..\MyFile.h" // relative path
```

Define

Another important directive is `#define`, which is used to create compile-time constants, also called *macros*. After this directive, the name of the constant is specified followed by what it will be replaced by.

```
#define PI 3.14 // macro definition
```

The preprocessor will go through and change any occurrences of this constant with whatever comes after it in its definition until the end of the line.

```
double d = PI; // d = 3.14
```

By convention, macros are named using uppercase letters with each word separated by an underscore. That way they are easy to spot when reading the source code.

Undefine

A `#define` directive should not be used to directly override a previously defined macro. Doing so will produce a compiler warning. In order to change a macro, it first needs to be undefined using the `#undef` directive. Attempting to undefine a macro that is not currently defined will not generate a warning.

```
#undef PI // undefine  
#undef PI // allowed
```

Predefined Macros

There are a number of macros that are predefined by the compiler. To distinguish them from user-defined macros, their names typically begin and end with two underscores. The following table lists some of the more useful predefined macros.

Directive	Description
<code>__FILE__</code>	Name and path of the current source file
<code>__LINE__</code>	Current line number
<code>__DATE__</code>	Compilation date in MMM DD YYYY format
<code>__TIME__</code>	Compilation time in HH:MM:SS format
<code>__func__</code>	Name of the current function; added in C++11

A common use for predefined macros is to provide debugging information. To give an example, the following error message includes the file name and line number where the message occurs.

```
cout << "Error in " << __FILE__ << " at line " << __LINE__;
```

Macro Functions

Macros can be made to take arguments. This allows them to define compile-time functions. For example, the following macro function gives the square of its argument.

```
#define SQUARE(x) ((x)*(x))
```

The macro function is called just as if it were a regular C++ function. Keep in mind that for this kind of function to work, the arguments must be known at compile time.

```
int x = SQUARE(2); // 4
```

Note the extra parentheses in the macro definition. They are used to avoid problems with operator precedence. Without the parentheses, the following example would give an incorrect result, as the multiplication would then be carried out before the addition.

```
#define SQUARE(x) x*x

int main()
{
    int x = SQUARE(1+1); // 1+1*1+1 = 3
}
```

To break a macro function across several lines, you use the backslash character. This will escape the newline character that marks the end of a preprocessor directive. For this to work, there must not be any whitespace after the backslash.

```
#define MAX(a,b) \
(a)>(b) ? \
(a): (b)
```

Although macros can be powerful, they tend to make the code more difficult to read and debug. Macros should therefore only be used when they are absolutely necessary and should always be kept short. C++ code—such as constant variables, enum classes, and constexpr functions—can often accomplish the same goal more efficiently and safely than #define directives can.

```
#define DEBUG 0
const bool debug = 0;

#define FORWARD 1
#define STOP 0
#define BACKWARD -1
enum class dir { forward = 1, stop = 0, backward = -1 };

#define MAX(a,b) (a)>(b) ? (a): (b)
constexpr int max(int a, int b) { return a>b ? a:b; }
```

Conditional Compilation

The directives used for conditional compilation can include or exclude part of the source code if a certain condition is met. First, there is the `#if` and `#endif` directives, which specify a section of code that will be included only if the condition after the `#if` directive is true. Note that this condition must evaluate to a constant expression.

```
#define DEBUG_LEVEL 3

#if DEBUG_LEVEL > 2
    // ...
#endif
```

Just as with the C++ `if` statement, any number of `#elif` (else if) directives and one final `#else` directive can be included.

```
#if DEBUG_LEVEL > 2
    // ...
#elif DEBUG_LEVEL == 2
    // ...
#else
    // ...
#endif
```

Conditional compilation also provides a useful means of temporarily commenting out large blocks of code for testing purposes. This often cannot be done with the regular multiline comment since they cannot be nested.

```
#if 0
    /* Removed from compilation */
#endif
```

Compile if Defined

Sometimes, a section of code should be compiled only if a certain macro has been defined, irrespective of its value. For this purpose, two special operators can be used: `defined` and `!defined` (not defined).

```
#define DEBUG

#if defined DEBUG
    // ...
#elif !defined DEBUG
    // ...
#endif
```

The same effect can also be achieved using the directives `#ifdef` and `#ifndef`, respectively. For instance, the `#ifdef` section is compiled only if the specified macro has been previously defined. Note that a macro is considered defined even if it has not been given a value.

```
#ifdef DEBUG
    // ...
#endif

#ifndef DEBUG
    // ...
#endif
```

Error

When the `#error` directive is encountered, the compilation is aborted. This directive can be useful to determine whether a certain line of code is being compiled. It can optionally take a parameter that specifies the description of the generated compilation error.

```
#error Compilation aborted
```


Line

A less commonly used directive is `#line`, which can change the line number that is displayed when an error occurs during compilation. Following this directive, the line number will as usual be increased by one for each successive line. The directive can take an optional string parameter that sets the file name that will be shown when an error occurs.

```
#line 5 "myapp.cpp"
```

Pragma

The last standard directive is `#pragma`, or pragmatic information. This directive is used to specify options to the compiler, and as such, they are vendor specific. To give an example, `#pragma message` can be used with many compilers to output a string to the build window. Another common argument for this directive is `warning`, which changes how the compiler handles warnings.

```
// Show compiler message
#pragma message("Hello Compiler")

// Disable warning 4507
#pragma warning(disable : 4507)
```

Attributes

A new standardized syntax was introduced in C++11 for providing compiler-specific information in the source code, so-called *attributes*. Attributes are placed within double square brackets and may, depending on the attribute, be applied to any code entities. To give an example, a

standard attribute added in C++14 is `[[deprecated]]`, which indicates that the use of a code entity has become discouraged.

```
// Mark as deprecated
[[deprecated]] void foo() {}
```

This attribute allows the compiler to emit a warning whenever such an entity is used. A message can be included in this warning to describe why the entity has been deprecated.

```
[[deprecated("foo() is unsafe, use bar() instead")]]
void foo() {}
```

Another example is the `[[noreturn]]` attribute, which specifies to the compiler that a function will not return to the calling function. This may be the case for functions that loop forever, throw exceptions, or exit the application.

```
[[noreturn]] void f()
{
    exit(0); // terminate program
}
```

The compiler may use this attribute for making optimizations as well as providing a warning that any statement following a call to this function will be unreachable.

CHAPTER 25

Exception Handling

Exception handling allows developers to deal with unexpected situations that may occur in a program.

Throwing Exceptions

When a function encounters a situation that it cannot recover from, it can generate an exception to signal the caller that the function has failed. This is done using the `throw` keyword followed by whatever it is the function wants to signal. When this statement is reached, the function will stop executing and the exception will propagate up to the caller where it can be caught, using a `try-catch` statement.

```
double divide(double x, double y)
{
    if (y == 0) throw 0;
    return x / y;
}
```

Try-Catch Statement

The `try-catch` statement consists of a `try` block containing code that may cause exceptions and one or more `catch` clauses to handle them. In the previous case, an integer is thrown so a `catch` block needs to be included that handles this type of exception. The thrown exception will

get passed as an argument to this exception handler, where it can be used to determine what has gone wrong with the function. Note that when the exception has been handled, the execution will then continue running after the try-catch blocks and not after the throw statement.

```
try {
    divide(10,0);
}
catch(const int& e) {
    cout << "Error code: " << e;
}
```

An exception handler can catch a thrown expression by value, by reference, or by pointer. However, catching by value should be avoided since this causes an extra copy to be made. Catching by const reference is generally preferable. If the code in the try block can throw more types of exceptions, then more catch clauses need to be added to handle them as well. Keep in mind that only the handler that matches the thrown expression will be executed, and the handlers are tried in the order they appear in the code.

```
catch(const char& e) {
    cout << "Error char: " << e;
}
```

To catch all types of exceptions, an ellipsis (...) can be used as the parameter of catch. This default handler must be placed as the last catch statement since no handler placed after it will ever be executed.

```
catch(...) { cout << "Error"; }
```

Rethrowing Exceptions

If an exception handler cannot recover from an exception, it can be rethrown by using the `throw` keyword with no argument specified. This will pass the exception up the call stack until another try-catch block is encountered. Be careful however, because if an exception is never caught, the program will terminate with a runtime error.

```
int main()
{
    try {
        try { throw 0; }
        catch(...) { throw; } // rethrow exception
    }
    catch(...) { throw; } // runtime error
}
```

Noexcept Specifier

The `noexcept` specifier indicates that a function is intended not to throw any exceptions. The main benefit of using `noexcept` is that it enables certain compiler optimizations, because the specifier allows the program to terminate without unwinding the call stack if for any reason an exception still occurs.

```
void foo() noexcept {} // must not throw exceptions
void bar() {} // may throw exceptions
```

Since C++11, the `noexcept` specifier may also be used as a compile-time operator to check if a function is declared to not throw any exceptions. Note that as of C++17, the exception specification has become a part of the type system, so the `noexcept` property needs to be included when binding a function pointer to such a function.

```
void(*pFunc)() noexcept = foo; // function pointer
pFunc(); // call function through pointer
cout << noexcept(pFunc); // "1" (true)
```

In this example, `pFunc` is a pointer to a function that takes zero arguments and returns `void`.

Exception Class

As previously mentioned, any data type can be thrown in C++. However, the standard library does provide a base class called `exception`, which is specifically designed to declare objects to be thrown. More specific exceptions can be created by deriving from this base class or from other exception classes available in the standard library. The `exception` class is defined in the `exception` header file and is located under the `std` namespace. As seen in the following code, the class can be constructed with a string that becomes the exception's description.

```
#include <exception>
using namespace std;

void makeError()
{
    throw exception("My Error Description");
}
```

When catching this exception, the object's function called `what` can be used to retrieve the description.

```
try {
    makeError();
}
catch (const exception& e) {
    cout << e.what(); // "My Error Description"
}
```

CHAPTER 26

Type Conversions

Converting an expression from one type to another is known as *type conversion*. This can be done implicitly or explicitly.

Implicit Conversions

An implicit conversion is performed automatically by the compiler when an expression needs to be converted into one of its compatible types. For example, any conversions between the primitive data types can be done implicitly.

```
long a = 5; // int implicitly converted to long
double b = a; // long implicitly converted to double
```

These implicit primitive conversions can be further grouped into two kinds: *promotion* and *demotion*. Promotion occurs when an expression gets implicitly converted into a larger type, and demotion occurs when converting an expression to a smaller type. Because a demotion can result in the loss of precision, these conversions will generate a warning on most compilers. If the potential information loss is intentional, the warning can be suppressed by using an explicit cast.

```
// Promotion
long a = 5; // int promoted to long
double b = a; // long promoted to double
```

```
// Demotion
int c = 10.5; // warning: possible loss of data
bool d = c; // warning: possible loss of data
```

Explicit Conversions

The first explicit cast is the one inherited from C, commonly called the *C-style cast*. The desired data type is simply placed in parentheses to the left of the expression that needs to be converted. This cast should be avoided in modern C++ code.

```
int c = (int)10.5; // double demoted to int
char d = (char)c; // int demoted to char
```

C++ Casts

The C-style cast is suitable for most conversions between the primitive data types. However, when it comes to conversions between objects and pointers, it can be too powerful. In order to get greater control over the different types of conversions possible, C++ introduced four new casts, called *named casts* or *new-style casts*. These casts are static, reinterpret, const, and dynamic cast.

```
static_cast<new_type> (expression)
reinterpret_cast<new_type> (expression)
const_cast<new_type> (expression)
dynamic_cast<new_type> (expression)
```

As seen here, their format is to include the cast's name with the new type enclosed in angle brackets followed by the expression to be converted in parentheses. These casts allow more precise control over how a conversion should be performed, which in turn makes it easier

for the compiler to catch conversion errors. In contrast, the C-style cast includes the static, reinterpret, and const cast in one operation. That cast is therefore more likely to execute subtle conversion errors if used incorrectly.

Static Cast

The static cast performs conversions between compatible types. It is similar to the C-style cast, but more restrictive. For example, the C-style cast would allow an integer pointer to point to a char.

```
char c = 10; // 1 byte
int *p = (int*)&c; // 4 bytes
```

Since this results in a four-byte pointer pointing to one byte of allocated memory, writing to this pointer will either cause a runtime error or overwrite some adjacent memory.

```
*p = 5; // runtime error: stack corruption
```

In contrast to the C-style cast, the static cast will allow the compiler to check that the pointer and pointee data types are compatible, which allows the programmer to catch this incorrect pointer assignment during compilation.

```
int *q = static_cast<int*>(&c); // compile-time error
```

Reinterpret Cast

To force the pointer conversion, in the same way as the C-style cast does in the background, the reinterpret cast would be used instead.

```
int *r = reinterpret_cast<int*>(&c); // forced conversion
```

This cast handles conversions between certain unrelated types, such as from one pointer type to another incompatible pointer type. It will simply perform a binary copy of the data without altering the underlying bit pattern. Note that the result of such a low-level operation is system specific and therefore not portable. It should be used with caution if it cannot be avoided altogether.

Const Cast

The third C++ cast is the const cast. This one is primarily used to add or remove the const modifier of a variable.

```
const int myConst = 5;
int *nonConst = const_cast<int*>(&myConst); // removes const
```

Although the const cast allows the value of a constant to be changed, doing so is still invalid code that may cause a runtime error. This could occur, for example, if the constant was located in a section of read-only memory.

```
*nonConst = 10; // potential runtime error
```

Const cast is instead used mainly when there is a function that takes a non-constant pointer argument, even though it does not modify the pointee.

```
void print(int *p) { std::cout << *p; }
```

The function can then be passed a constant variable by using a const cast.

```
print(&myConst); // error: cannot convert const int* to int*
print(nonConst); // allowed
```

C-Style and New-Style Casts

Keep in mind that the C-style cast can also remove the `const` modifier, but again since it does this conversion behind the scenes, the C++ casts are preferable. Another reason to use the C++ casts is that they are easier to find in the source code than the C-style casts. This is important because casting errors can be difficult to discover. A third reason for using the C++ casts is that they are unpleasant to write. Since explicit conversions in many cases can be avoided, this was done intentionally so that programmers would look for a different solution.

Dynamic Cast

The fourth and final C++ cast is the dynamic cast. This cast is only used to convert object pointers and object references into other pointers or reference types in the inheritance hierarchy. It is the only cast that makes sure that the object pointed to can be converted, by performing a runtime check that the pointer refers to a complete object of the destination type. For this runtime check to be possible, the object must be *polymorphic*. That is, the class must define or inherit at least one virtual function. This is because the compiler will only generate the needed runtime type information for such objects.

In the following code segment, a `MyChild` pointer is converted into a `MyBase` pointer using a dynamic cast. This derived-to-base conversion succeeds, because the `Child` object includes a complete `Base` object.

```
class MyBase { public: virtual void test() {} };
class MyChild : public MyBase {};

int main()
{
    MyChild *child = new MyChild();
```

```

    MyBase *base = dynamic_cast<MyBase*>(child); // ok
    // ...
    delete child;
}

```

The next example attempts to convert a `MyBase` pointer into a `MyChild` pointer. Since the `MyBase` object does not contain a complete `MyChild` object, this pointer conversion will fail. To indicate this, the dynamic cast returns a null pointer. This gives programmers a convenient way to check whether a conversion has succeeded during runtime.

```

MyBase *base = new MyBase();
MyChild *child = dynamic_cast<MyChild*>(base);
if (child == nullptr) cout << "Null pointer returned";
delete base;

```

If a reference is converted instead of a pointer, the dynamic cast will then fail by throwing a `bad_cast` exception. This needs to be handled using a try-catch statement.

```

#include <exception>
#include <iostream>
using namespace std;

class MyBase { public: virtual void test() {} };
class MyChild : public MyBase {};

int main()
{
    MyBase *base = new MyBase();
    try {
        MyChild &child = dynamic_cast<MyChild&>(*base);
    }
}

```

```

    catch(const bad_cast &e) {
        cout << e.what(); // "bad dynamic_cast"
    }
    delete base;
}

```

Dynamic or Static Cast

The advantage of using a dynamic cast is that it allows the programmer to check whether a conversion has succeeded during runtime. The disadvantage is that there is a performance overhead associated with doing this check. For this reason, using a static cast would have been preferable in the first example, because a derived-to-base conversion will never fail.

```
MyBase *base = static_cast<MyBase*>(child); // ok
```

However, in the second example, the conversion may either succeed or fail. It will fail if the `MyBase` object contains a `MyBase` instance, and it will succeed if it contains a `MyChild` instance. In some situations, this may not be known until runtime. When this is the case, a dynamic cast is a better choice than a static cast.

```
// Succeeds for a MyChild object
MyChild *child = dynamic_cast<MyChild*>(base);
```

If the base-to-derived conversion had been performed using a static cast instead of a dynamic cast, the conversion would not have failed. It would have returned a pointer that referred to an incomplete object. Dereferencing such a pointer can lead to runtime errors.

```
// Allowed, but invalid
MyChild *child = static_cast<MyChild*>(base);
```

CHAPTER 27

Smart Pointers

Several smart pointer classes were added in C++11 for managing dynamically allocated memory and resources in general. By using these container classes, instead of raw pointers, it is no longer necessary to manually delete objects created with the `new` keyword. This simplifies coding by helping to prevent memory leaks.

Unique Pointer

The first smart pointer that we look at is the unique pointer (`std::unique_ptr`), which simply acts as a container for a raw pointer. It replaces another deprecated smart pointer named `auto_ptr`, which was removed in C++17. Consider the following example on how to use a unique pointer.

```
#include <memory> // include smart pointers
#include <iostream>
using namespace std;

struct Foo
{
    int val;
    Foo() { cout << "1"; }
    ~Foo() { cout << "3"; }
};
```

```
int main()
{
    unique_ptr<Foo> p(new Foo()); // "1"
    p->val = 2;
    cout << p->val; // "2"
} // "3"
```

The output of this code is "123" as the pointer is created, used, and then destroyed automatically when it goes out of scope. Note that the smart pointer is created not through assignment but instead by passing a raw pointer to its constructor. Once created, however, the smart pointer is used just as a regular pointer, in this case with the arrow operator (`->`) to dereference the pointer and access the member of the object in a single operation.

As the name implies, a unique pointer has exclusive ownership of the object it points to and therefore cannot be copied. It can, however, transfer ownership to another unique pointer using the `std::move` function. After completing such a transfer, the original pointer will automatically be set to `nullptr`.

```
unique_ptr<Foo> u1(new Foo());
unique_ptr<Foo> u2 = u1; // compile-time error
unique_ptr<Foo> u3 = move(u1); // transfers ownership
```

Shared Pointer

In cases where shared ownership of a dynamically allocated object is necessary, there is the shared pointer (`std::shared_ptr`). Unlike the unique pointer, a shared pointer can be copied. The memory to the object will not be deallocated until the last remaining shared pointer owning the object is destroyed, either by going out of scope or by resetting the pointer to `nullptr` manually.

```
shared_ptr<Foo> s1(new Foo());
shared_ptr<Foo> s2 = s1; // extends ownership
s1 = nullptr; // reset pointer
s2 = nullptr; // reset last pointer and delete memory
```

As of C++14, the use of the `new` keyword is discouraged in most circumstances. Instead, the `std::make_unique` and `std::make_shared` functions are recommended when allocating dynamic memory.

```
unique_ptr<Foo> u = make_unique<Foo>();
shared_ptr<int> s = make_shared<int>(10);
```

Both of these helper methods perform value initialization. Since C++20 there are also methods available for doing default initialization. This avoids unnecessary initialization in situations where the initial value is not needed. Type deduction is used here to avoid having to type the type twice.

```
auto u2 = make_unique_for_overwrite<Foo>();
auto s2 = make_shared_for_overwrite<int>(10);
```

Weak Shared Pointer

A weak shared pointer (`std::weak_ptr`) can be created from a shared pointer. Unlike the shared pointer, a weak shared pointer is non-owning, meaning that the object will be cleaned up when all shared pointers go out of scope, regardless of any weak shared pointers. In order to access the referenced object, a weak shared pointer must first be converted into a shared pointer using the `lock` method. Here is an example to illustrate.

```
#include <memory>
#include <iostream>
using namespace std;

void observe(weak_ptr<int> weak)
```



```
{
    shared_ptr<int> s = weak.lock();
    if (s != nullptr) {
        cout << "Pointer is " << *s << endl;
    }
    else {
        cout << "Pointer has expired" << endl;
    }
}

int main()
{
    shared_ptr<int> s = make_shared<int>(5);
    weak_ptr<int> w = s; // copy pointer without ownership
    observe(w); // "Pointer is 5"
    s = nullptr; // delete managed object
    observe(w); // "Pointer has expired"
}
```

CHAPTER 28

Templates

Templates provide a way to make a class, function, or variable operate with different data types without having to rewrite the code for each type.

Function Templates

This example shows a function that swaps two integer arguments.

```
void swap(int& a, int& b)
{
    int tmp = a;
    a = b;
    b = tmp;
}
```

To convert this method into a function template that can work with any type, the first step is to add a *template parameter declaration* before the function. This declaration includes the `template` keyword followed by the keyword `typename` and the name of the *template type parameter*, both enclosed between angle brackets. The name of the template parameter may be anything, but it is common to name it with a capital `T`.

```
template<typename T>
```

Alternatively, the keyword `class` can be used instead of `typename`. They are equivalent in this context.

```
template<class T>
```

The second step in creating a function template is to replace the data type that will be made generic with the template type parameter.

```
template<class T>
void swap(T& a, T& b)
{
    T tmp = a;
    a = b;
    b = tmp;
}
```

Calling Function Templates

The function template is now complete. To use it, you can call `swap` as if it were a regular function, but with the desired template argument specified in angle brackets before the function arguments. Behind the scenes, the compiler will instantiate a new function with this template parameter filled in, and it is this generated function that will be called from this line.

```
int a = 1, b = 2;
swap<int>(a,b); // calls int version of swap
```

Every time the function template is called with a new type, the compiler will instantiate another function using the template.

```
bool c = true, d = false;
swap<bool>(c,d); // calls bool version of swap
```

In this example, the swap function template may also be called without specifying the template parameter. This is because the compiler can automatically determine the type, because the function template's arguments use the template type. However, if this is not the case, or if there is a need to force the compiler to select a specific instantiation of the function template, the template parameter would then need to be explicitly specified within angle brackets.

```
int e = 1, f = 2;
swap(e,f); // calls int version of swap
```

Multiple Template Parameters

Templates can be defined to accept more than one template parameter by adding them between the angle brackets separated by commas.

```
template<class T, class U>
void swap(T& a, U& b)
{
    T tmp = a;
    a = b;
    b = tmp;
}
```

The second template parameter in this example allows swap to be called with two arguments of different types.

```
int main()
{
    int a = 1;
    long b = 2;
    swap<int, long>(a,b);
    swap(a,b); // alternative
}
```

Class Templates

Class templates allow class members to use template parameters as types. They are created in the same way as function templates.

```
template<class T>
class MyBox
{
public:
    T a, b;
    MyBox(const T& x, const T& y) : a(x), b(y) {}
};
```

The compiler can deduce the template type parameters if they are based on the arguments passed to a constructor of the class.

```
int main()
{
    // Without type deduction
    MyBox<int> box(1, 2); // MyBox<int>

    // With type deduction
    MyBox box(2.1, 3.2); // MyBox<double>
}
```

Another point to remember when using class templates is that if a method is defined outside of the class template, that definition must also be preceded by the template declaration.

```
template<class T>
class MyBox
{
public:
    T a, b;
    void swap();
};
```

```
template<class T>
void MyBox<T>::swap()
{
    T tmp = a;
    a = b;
    b = tmp;
}
```

Notice that the template parameter is included in the swap template function definition after the class name qualifier. This specifies that the function's template parameter is the same as the template parameter of the class.

Non-type Parameters

In addition to type parameters, both class and function templates can also have regular function-like parameters. As an example, the unsigned int template parameter is used to specify the size of an array.

```
template<class T, unsigned int N>
class MyBox
{
public:
    T store[N];
};
```

When this class template is instantiated, both a type and an integer have to be included.

```
MyBox<int, 5> box;
```

Default Types and Values

Class and function template parameters can be given default values and types.

```
template<class T = int, int N = 5>
```

To use these defaults, the angle brackets just need to be left empty when instantiating the class template.

```
MyBox<> box;
```

Class Template Specialization

If there is a need to define a different implementation for a template when a specific type is passed as the template parameter, a *template specialization* can be declared. For example, in the following class template, there is a `print` method that outputs the value of a class template field.

```
#include <iostream>

template<class T>
class MyBox
{
public:
    T a;
    void print() { std::cout << a; }
};
```

When the template parameter is a `bool`, the method should print out “true” or “false” instead of “1” or “0”. One way to do this is to create a *class template specialization*. A reimplementaion of the class template is then created where the template parameter list is empty. Instead, a `bool`

specialization parameter is placed after the class template's name, and this data type is used instead of the template parameter throughout the implementation.

```
template<>
class MyBox<bool>
{
public:
    bool a;
    void print() { std::cout << (a ? "true" : "false"); }
};
```

When this class template is instantiated with a `bool` template type, this template specialization will be used instead of the standard one.

```
int main()
{
    MyBox<bool> box { true };
    box.print(); // "true"
}
```

Note that there is no inheritance of members from the standard template to the specialized template. The whole class will have to be redefined.

Function Template Specialization

Since there is only one function that is different between the templates in the previous example, a better alternative is to create a *function template specialization*. This kind of specialization looks very similar to the class template specialization, but is only applied to a single function instead of the whole class.


```

#include <iostream>

template<class T>
class MyBox
{
public:
    T a;
    template<class T> void print() {
        std::cout << a;
    }
    template<> void print<bool>() {
        std::cout << (a ? "true" : "false");
    }
};

```

This way, only the print method has to be redefined and not the whole class.

```

int main()
{
    MyBox<bool> box = { true };
    box.print<bool>(); // "true"
}

```

Notice that the template parameter has to be specified when the specialized function is invoked. This is not the case with the class template specialization.

Variable Templates

In addition to function and class templates, C++14 allows variables to be templated. This is achieved using the regular template syntax.

```

template<class T>
constexpr T pi = T(3.1415926535897932384626433L);

```

Together with the `constexpr` specifier, this template allows the value of the variable to be computed at compile time for a given type, without having to type cast the value.

```
int i = pi<int>; // 3
float f = pi<float>; // 3.14...
```

Variadic Templates

C++11 allows template definitions to take a variable number of type arguments. To illustrate, consider the following function, which returns the sum of any number of ints passed to it.

```
#include <iostream>
#include <initializer_list>
using namespace std;

int sum(initializer_list<int> numbers)
{
    int total = 0;
    for(auto& i : numbers) { total += i; }
    return total;
}
```

The `initializer_list` type indicates that the function accepts a brace-enclosed list as its argument, so the function must be called in this manner.

```
int main()
{
    cout << sum( { 1, 2, 3 } ); // "6"
}
```

The next example changes this function into a variadic template function. Such a function is traversed recursively rather than iteratively, so once the first argument has been handled, the function calls itself with the remaining arguments.

The variadic template parameter is specified using the ellipsis (...) operator, followed by a name. This defines a so-called parameter pack. The parameter pack is bound to a parameter in the function (... rest) and then unpacked into separate arguments (rest ...) when the function calls itself recursively.

```
int sum() { return 0; } // end condition

template<class T0, class ... Ts>
decltype(auto) sum(T0 first, Ts ... rest)
{
    return first + sum(rest ...);
}
```

This variadic template function can be called as a regular function, with any number of arguments. In contrast to the previously defined variadic function, this template function accepts arguments of any type.

```
int main()
{
    cout << sum(1, 1.5, true); // "3.5"
}
```

Fold Expressions

C++17 introduced fold expressions, which make it possible to apply a binary operator to all elements of a parameter pack in one statement. This allows the previous variadic template function to be written more concisely and without the use of recursion.

```
template<class... T>
decltype(auto) sum(T... args)
{
    // Unpacks to: a1 + (a2 + (a3 + a4))...
    return (args + ...);
}
```

A unary right fold is here performed in the return statement, expanding the parameter pack starting from the left and applying the binary operator to all arguments before returning the result. Parameter packs may also be unpacked from right to left, by placing the ellipsis to the left of the parameter pack, as shown in the following example using the subtraction operator.

```
#include <iostream>
using namespace std;

template<class... T>
decltype(auto) difference(T... args)
{
    // Unpacks to: ...(a1 - a2) - a3
    return (... - args);
}

int main()
{
    cout << difference(5, 2, 1); // "2" (5-2-1)
}
```

Concepts

A concept is a named set of constraints that limit what template arguments may be used with a template. They were introduced in C++20 to allow template arguments to be type-checked at compile time. The following example defines a concept named `MyIntegral` which requires the type to be convertible to a whole number type. The `is_integral_v` class template used here is part of the standard library, and it is evaluated as true if `T` is an integral type.

```
#include <concepts>
#include <type_traits>
// Concept declaration
template <class T>
concept MyIntegral = std::is_integral_v<T>;
```

This concept can be applied to constrain template arguments, such as for the following function template. Any template argument used to initialize this function template must satisfy the requirement of the concept, or else the compilation will fail.

```
template<MyIntegral T>
bool is_positive(T a)
{
    return a > 0;
}

int main()
{
    is_positive(5); // ok, int satisfies MyIntegral
    is_positive("Hi"); // error, string does not satisfy MyIntegral
}
```

The standard library includes a number of predefined concepts that should be used in favor of user-defined ones whenever possible. In this example, the standard concept `std::integral` performs the same function as “MyIntegral,” so the preceding function template can be redefined as follows.

```
#include <concepts>
template<std::integral T>
bool is_positive(T a)
{
    return a > 0;
}
```

There are two ways to express a concept. The first way is in the form of a conditional expression, which was the form used for the `integral` concept defined earlier. The following example makes use of the `integral` concept and also adds a second constraint to make sure the type is signed and not unsigned. Note that this constraint makes use of the fact that constructing an unsigned type with a negative value returns a positive value, because the unsigned type cannot represent the negative value.

```
template <class T>
concept Signed_Integral = std::integral<T> && T{-1} < T{0};
```

The second way to define a concept is to use a `requires` clause. This clause defines objects of the types to be tested and then a list of one or more constraints. Each constraint consists of an expression in curly brackets followed by the expected return type. If all constraints are true, the compiler will allow the type. For instance, the following concept declares that the type must implement both the equal to and not equal to operators and that the result of these operations must be convertible to a `bool`.

```

template<class T>
concept Equal = requires(T a, T b)
{
    { a == b } -> bool;
    { a != b } -> bool;
};

template<Equal T>
bool areEqual(T x, T y)
{
    return x == y;
}

int main()
{
    areEqual(1, 1); // true
}

```

Abbreviated Function Templates

Function templates can be abbreviated as of C++20 by using the `auto` placeholder type. When `auto` appears in the parameter list, the function automatically becomes a function template and the `auto` parameter becomes its template parameter. Applying a concept to such a function is done by adding the name of the concept before the type in the parameter list. Bear in mind that abbreviated function templates are not supported in Visual Studio 2019 as of version 16.3.

```

#include <concepts>
bool is_positive(std::integral auto a)
{
    return a > 0;
}

```

```
int main()
{
    is_positive(2); // calls int version
    is_positive(3L); // calls long version
}
```

Template Lambdas

With C++14 generic lambdas were introduced, which meant that parameters declared as `auto` became template parameters. The following example defines a generic lambda that returns the size of a vector.

```
#include <iostream>
#include <vector>
using namespace std;
int main()
{
    vector<int> v { 1, 2, 3 };
    auto get_size = [](const auto& v) { return size(v); };
    cout << get_size(v); // 3
}
```

It would be preferable to restrict this lambda to only work with vector types. This ability was added in C++20 by allowing full use of template parameters in lambdas.

```
auto get_size = []<typename T>(vector<T> const& v) { return
size(v); };
```


CHAPTER 29

Headers

When a project grows, it is common to split the code up into different source files. When this happens, the interface and implementation are generally separated. The interface is placed in a header file, which commonly has the same name as the source file and an `.h` file extension. This header file contains forward declarations for the source file entities that need to be accessible to other compilation units in the project. A compilation unit consists of a source file (`.cpp`) and any included header files (`.h` or `.hpp`).

Why Use Headers

C++ requires everything to be declared before it can be used. It is not enough to simply compile the source files in the same project. For example, if a function is placed in `MyFunc.cpp`, and a second file named `MyApp.cpp` in the same project tries to call it, the compiler will report that it cannot find the function.

```
// MyFunc.cpp
void myFunc() {}

// MyApp.cpp
int main()
{
    myFunc(); // error: myFunc identifier not found
}
```

To make this work, the function's prototype has to be included in `MyApp.cpp`.

```
// MyApp.cpp
void myFunc(); // prototype

int main()
{
    myFunc(); // ok
}
```

Using Headers

This can be made more convenient if the prototype is placed in a header file named `MyFunc.h`, and this header is included in `MyApp.cpp` through the use of the `#include` directive. This way if any changes are made to `MyFunc.cpp` or `MyFunc.h`, there is no need to update the prototypes in `MyApp.cpp`. Furthermore, any source file that wants to use the shared code in `MyFunc.cpp` can just include this one header.

```
// MyFunc.h
void myFunc(); // prototype

// MyApp.cpp
#include "MyFunc.h"
```

What to Include in Headers

As far as the compiler is concerned, there is no difference between a header file and a source file. The distinction is only conceptual. The key idea is that the header should contain the interface of the implementation file—that is, the code that other source files will need to use. This may include, for instance, shared constants, macros, and type aliases. Headers

should not include using namespace directives, because that would force the namespace inclusion upon everyone using the header.

```
// MyApp.h - Interface
#define DEBUG 0
const double E = 2.72;
using ulong = unsigned long;
```

As already mentioned, the header can contain prototypes of the shared functions defined in the source file.

```
void myFunc(); // prototype
```

Additionally, shared classes are typically specified in the header, while their methods are implemented in the source file.

```
// MyApp.h
class MyClass
{
public:
    void myMethod();
};

// MyApp.cpp
void MyClass::myMethod() {}
```

As with functions, it is necessary to forward declare global variables before they can be referenced in a compilation unit outside the one containing their definition. This is done by placing the shared variable in the header and marking it with the keyword `extern`. This keyword indicates that the variable is initialized in another compilation unit. Functions are `extern` by default, so function prototypes do not need to include this specifier. Keep in mind that global variables and functions may be declared externally multiple times in a program, but they may be defined only once.

```
// MyApp.h
extern int myGlobal;

// MyApp.cpp
int myGlobal = 0;
```

It should be noted that the use of shared global variables is discouraged. This is because the larger a program becomes, the more difficult it is to keep track of which functions access and modify these variables. The preferred method is to instead pass variables to functions only as needed, in order to minimize the scope of those variables.

The header should not include any executable statements, with two exceptions. First, if a shared class method or global function is declared as `inline`, that function must be defined in the header. Otherwise, calling the inline function from another source file will give an unresolved external error. Note that the `inline` modifier suppresses the single definition rule that normally applies to code entities.

```
// MyApp.h
inline void inlineFunc() {}

class MyClass
{
public:
    void inlineMethod() {}
};
```

The second exception is shared templates. When encountering a template instantiation, the compiler needs to have access to the implementation of that template, in order to create an instance of it with the type arguments filled in. The declaration and implementation of templates are therefore generally put into the header file all together.

```
// MyApp.h
template<class T>
class MyTemp { /* ... */ };

// MyApp.cpp
MyTemp<int> o;
```

Instantiating a template with the same type in many compilation units leads to significant redundant work done by the compiler and linker. To prevent this, C++11 introduced extern template declarations. A template instantiation marked as extern signals to the compiler not to instantiate the template in this compilation unit.

```
// MyApp.cpp
MyTemp<int> b; // instantiation is done here

// MyFunc.cpp
extern MyTemp<int> a; // suppress redundant instantiation
```

If a header requires other headers, it is common to include those files as well, to make the header stand-alone. This ensures that everything needed is included in the correct order, solving potential dependency problems for every source file that requires the header.

```
// MyApp.h
#include <stddef.h> // include size_t
void mySize(std::size_t);
```

Note that since headers mainly contain declarations, any extra headers included should not affect the size of the program, although they may slow down compilation.

Inline Variables

As of C++17, variables may be specified as inline, in addition to functions and methods. This allows constant and static variables to be defined in a header file, because the inline modifier removes the single definition rule that would normally prevent this. Once an inline variable has been defined, all compilation units referencing that header will use the same definition.

```
// MyApp.h
struct MyStruct
{
    static const int a;
    inline static const int b = 10; // alternative
};
inline int const MyStruct::a = 10;
```

The `constexpr` keyword implies inline, so a variable declared as `constexpr` may also be initialized in a header file. However, such a variable must be initialized to a compile-time constant.

```
struct MyStruct {
    static constexpr int a = 10;
};
```

An inline variable is not restricted to only constant expressions, as seen in the following example where the inline variable is initialized to a random value between 1 and 6. This value is guaranteed to be the same for all compilation units using this header, even though the value is not set until runtime.

```
#include <cstdlib> // rand, srand
#include <ctime> // time

struct MyStruct {
    static const int die;
};

inline const int MyStruct::die =
    (srand((unsigned)time(0)), rand()%6+1); // 1-6
```

Note the use of the comma operator here, which evaluates the left expression first and then evaluates and returns the right expression. The left expression uses the current time to seed the random number generator with the `srand` function. The right expression retrieves a random integer with the `rand` function and formats the integer into the 1–6 range.

Include Guards

An important point to bear in mind when using header files is that a shared code entity may only be defined once. Consequently, including the same header file more than once will likely result in compilation errors. A common way to prevent this is to use a so-called *include guard*. An include guard is created by enclosing the whole header in an `#ifndef` section that checks for a macro specific to that header file. Only when the macro is not defined is the file included. The macro is then defined, which effectively prevents the file from being included again.

```
// MyApp.h
#ifndef MYAPP_H
#define MYAPP_H
// ...
#endif // MYAPP_H
```

Most compilers also support the nonstandard `#pragma once` directive, which serves the same purpose as include guards but with less code. Just place the directive in the header file to make sure it can only be included once.

```
#pragma once
```

Before including a header file, it may also be a good idea to check if it exists. For this purpose, C++17 added the `__has_include` preprocessor expression, which evaluates to true if the header file is found.

```
#if __has_include("myapp.h")
#include "myapp.h"
#endif
```

Modules

A module is a set of one or more source code files that are compiled independently and can then be imported into another compilation unit. They were introduced in C++20 to remove common issues associated with using header files, such as header-order dependencies, naming collisions, and multiple inclusions of the same header file. Furthermore, as modules only need to be compiled once, they can reduce compilation times, especially for large projects.

To enable experimental support for modules in Visual Studio 2019 (version 16.3), right-click the project in the Solution Explorer and choose Properties. From there, select All Configurations from the Configuration drop-down list and then enable module support under Configuration Properties ► C/C++ ► Language ► Enable C++ Modules (experimental). Next, add a new file called `ModInterface.ixx` to the project by right-clicking the Source Files folder in the Solution Explorer and going to Add ► New item.

The `ixx` file extension is required for module interface units in Visual Studio. Some other compilers, such as GCC (GNU Compiler Collection), use a `cppm` file extension instead. An `export` module declaration is placed in the file to specify the name of the module.

```
// ModInterface.ixx
export module mymodule; // declare module name
```

Only code entities explicitly marked with `export` will be visible to source files using the module, which in the following example is the `getValue` function. All other code entities will be internal to the module and will not influence source code outside the module. This is a big advantage compared with header files, as headers may include code that inadvertently affect other parts of the code.

```
// ModInterface.ixx
export module mymodule;
#define VALUE 5
int hidden() { return VALUE; }
export int getValue() { return hidden(); }
```

Optionally, the implementation of the module can be separated from the interface unit into one or more module implementation units. Such an implementation file cannot export any names. Any code entities it declares will be visible across the entire module, but not outside the module to which it belongs. The implementation file itself may use any file extension.

```
// ModInterface.ixx
export module mymodule;
export int getValue();

// ModImplementation.cpp
module mymodule; // unit belongs to mymodule
```

```
#define VALUE 5
int hidden() { return VALUE; }
int getValue() { return hidden(); }
```

To get the `ixx` file to compile, right-click the file in the Solution Explorer and click Properties. From the Properties window, change the Item Type of the file to C/C++ Compiler and click OK. You will then be able to manually compile the file by right-clicking it and selecting Compile in the Solution Explorer.

With the module ready and compiled, it can be imported into any source file to make use of its functionality. The import declaration must appear at the global scope of the file importing the module.

```
// MyApp.cpp
import mymodule; // import module
#include <iostream>
using namespace std;

int main()
{
    cout << getValue() << endl; // "5"
}
```

Some standard library headers, such as `iostream` and `vector`, can be imported as if they were modules. This is not supported in Visual Studio 2019 (version 16.3). Keep in mind that an import declaration ends with a semicolon, unlike the include directive.

```
import <iostream>;
import <vector>;
```

Index

A

Access levels

- class declaration, 99
- friend classes and
 - functions, 102–104
- global function, 104
- guidelines, 101
- inheritance (public, protected/
private), 104
- private members, 99
- protected member, 100
- public access, 101

Aggregate

- initialization, 83, 116

Arithmetic operators, 17

Arrays, 33

- assignment, 33
- declaration/
 - allocation, 33
- delete keyword, 36
- dynamic memory, 35
- multi/single-dimensional
 - arrays, 34
- std::size function, 35–36
- vector, 36–37

Assignment operator

- (+=), 18, 39

B

Base class scoping, 96

Binary operator

- overloading, 122

Bitwise operator (|), 20

Booleantype, 16

C

Classes

- access object
 - members, 73–74
- definition, 71
- forward declaration, 74
- inline keyword, 72
- instance, 72
- methods, 71
- object creation, 72

Comments, 6

Comparison operators

- (<, >, <= / >=), 19, 123–126

Concatenation operator (+), 39

Conditional statements

- if statement, 45
- initializers, 48
- switch statement, 46
- ternary operator, 47

INDEX

Console compilation, 5

Constants

expressions

compilation error, 143

conditional statements, 145

functions and class

constructors, 143

integer and enumeration

types, 143

lambda expression, 144

object, 144

return value, 144

runtime evaluation, 144

static_assert declaration, 146

virtual functions, 145

fields, 142

guideline, 146

immediate function, 146

methods, 141

objects, 140

pointers, 139–140

references, 140

return type and method

parameters, 142

variables, 139

Constructors

aggregate initialization, 83

copy, 82

declaration, 75

default parameter, 78

default values, 75

designated initializers, 85

destructor, 79

direct initialization, 81

field initialization, 77–78

inheritance, 89–90

instance methods, 77

new initialization, 82

object initialization, 80

overloading, 76

special member functions, 79

this keyword, 77

uniform initialization, 83–84

value object, 81

Conversions

explicit constructor, 129

implicit constructor, 128–129

MyNum class, 127

operators

explicit constructors, 130–131

object declaration, 129–130

Copy initialization, 82

D

Data types/primitives, 7–8

Decltype keywords, 61–63

Designated initializers, 85–86

Direct initialization, 81

Downcasting, 88

Dynamic arrays, 35

E

Enumeration (enum)

constants, 109

constant values, 110

integer type, 113

- scope, 111
- switch statement, 109
- weakly typed enum, 111–113

Escape characters, 40–41

Exception handling, 157

- exception class, 160
- noexcept specifier, 159
- rethrown, 159
- throw keyword, 157
- try-catch statement, 157–158

F

Floating-point types, 13

Fold expressions, 182

Functions

- auto keyword, 61–63
- calling, 53
- capture clause, 67
- decltype, 62–64
- default values, 55
- definition, 53
- forward declaration, 56–57
- inline function, 60
- lambda function, 66–69
- multiple values, 63–65
- overloading, 55
- parameters, 54
- pass by address, 58
- pass by reference, 58
- pass by value, 57
- prototype, 56
- returns by value/reference/
address, 59

- return statement, 55–56
- structured bindings, 65

G

Generic lambdas, 187

GNU Compiler Collection
(GCC), 5, 197

H

Header files (.h/.hpp), 189

- constexpr keyword, 194
- getValue function, 197
- global variables and
functions, 191
- guards, 195–196
- inline variables, 194–195
- modules, 196–198
- namespace directives, 191
- shared functions/
templates, 191–192
- template instantiation, 193

Hello World

- header, resource/source files, 2
- integrated development
environment, 1
- IntelliSense, 4
- language standard selection, 2
- project creation, 2
- scope resolution operator, 3
- source file, 3–4
- standard namespace, 4

Hiding derived members, 93–94

I, J, K

Increment (++)/decrement (--)
operators, 18

Inequality operator (!=), 125

Inheritance

- access levels, 104
- constructor, 89–90
- downcasting, 88
- multiple inheritance, 91
- public keyword, 87
- upcasting, 87–88

Instantiate, *see* Constructors

int x = 1, y = 2, z, 9

Integrated Development

Environment (IDE), 1

L

Lambda function

- capture clause, 67–68
- decltype specifier, 69
- generic lambda
 - expressions, 66
- int arguments/returns, 66
- parameter list and return
 - type, 66–67
- stateless, 68

Logical operators (&&, ||/!), 19

Looping structures

- break/continue keyword, 52
- do-while, 50
- for, 50
- goto statement, 52
- while loop, 49

M

Multiple inheritance, 91

N

Namespaces

- access members, 134
- aliases, 136
- classes and functions, 133
- importing, 135
- member import, 135
- nesting, 134
- prototypes, 137
- scopes, 133
- type alias, 136

Nested namespace members, 134

NetBeans/Eclipse CDT, 1

NULL pointer, 26–28, 30

O

Octal/hexadecimal notation, 12

Noexcept specifier, 159

Operator overloading

- binary, 122
- comparison operator, 123–126
- MyNum objects and returns, 121
- overloadable operators, 126
- postfix operators, 123
- unary operators, 122–123
- user-defined types, 121

Operators

- increment (++)/
 - decrement (--), 18

- assignment(=), 18
- bitwise (|), 20
- arithmetic operators, 17
- comparison, 19
- logical operators, 19
- precedence, 20–21
- types, 17
- Overriding
 - base class scoping, 96
 - derived classes, 94–96
 - getArea method, 94
 - hiding derived members, 93–94
 - polymorphism, 94
 - pure virtual functions, 96

P, Q

- Prefix/postfix
 - operators, 123

Pointers

- asterisk (*), 23
- dereference operator (*), 24
- constants, 139, 140
- dynamic allocation, 25–26
- NULL constant, 26–28
- pass by address, 58
- pointee, 23
- point-pointer, 25
- references, 30
- static allocation, 25

Preprocessor

- attributes, 155
- conditional compilation, 152
- directives, 147

- error message/directive, 150, 153
- defined/!defined
 - (#ifdef/#ifndef), 153
- line number, 154
- macro functions, 148, 150–151
- pragma, 154
- predefined macros, 149
- source file, 148
- undefine directive, 149

R

References

- declaration, 29
- pointer, 30
- rvalue, 30–31

- Rethrown exceptions, 159

S

- Signed and unsigned integers, 12–13

- Single-line comment, 6

Smart pointers

- classes, 169
- pointers (*see* Pointers)
- shared pointer
 - (std::shared_ptr), 170–171
- unique pointer
 - (std::unique_ptr), 169–170
- weak shared pointer
 - (std::weak_ptr), 171–172

- Source file (.cpp), *see* Header files (.h/.hpp)

- Static keyword

INDEX

Static keyword (*cont.*)

- field (class field), 105–106

- global variables, 107

- local variables, 107

- methods, 106–107

std::get function, 64

std::make_tuple function, 64–65

std::tie function, 64–65

Strings

- combining string, 39–40

- compare option, 41

- declaration, 39

- direct/uniform initialization, 39

- encoding, 42

- escape characters, 40–41

- length and size functions, 41

- size_t/cstdint header, 41–42

- substr (substring), 42

- std::format function, 43

Structs

- aggregate initialization, 116

- anonymous, 116

- data structure, 115

- declaration syntax, 115

- initialization, 115–116

- uniform initialization, 116

T

Templates, 173

- abbreviated function

 - templates, 186

- calling function, 174

- class templates, 176–177

- concepts, 184–186

- default values and types, 178

- fold expressions, 182

- function, 173

- generic lambdas, 187

- multiple parameters, 175

- non-type parameters, 177

- parameter declaration, 173

- specialization

 - class template, 178–179

 - function, 179–180

- swap function, 175

- type parameters, 174

- variables, 180–181

- variadic templates, 181–182

Ternary operator, 47

Throwing exceptions, 157

Try-catch statement, 157–158

Type conversions

- casts, 162

- const cast, 164

- C-style/new-style casts, 165

- dynamic cast, 165–167

- explicit cast, 162

- implicit conversion, 161

- named cast/new-style

 - casts, 162

- polymorphic, 165

- promotion and

 - demotion, 161

- reinterpret cast, 163

- static cast, 163

- static/dynamic cast, 167

- try-catch statement, 166

U

Unary operators, [122–123](#)
Unicode characters, [43](#)
Uniform initialization, [83–84](#)
Union type
 anonymous, [118](#)
 benefit of, [117](#)
 integer field, [117–118](#)
Unique pointer (`std::unique_ptr`),
 [169–170](#)

V

Value initialization, [81](#)
Variables
 assignment operator (`=`), [8](#)
 bool type, [16](#)
 char type, [14](#)

data types, [7](#)
declaration, [8](#)
direct/uniform initialization, [9](#)
floating-point types, [13](#)
integer types, [10–11](#)
literal suffixes, [14](#)
numeric literals, [12](#)
scope of, [9–10](#)
signed and unsigned
 keywords, [12–13](#)
Variadic template, [181–182](#)
Vectors, [36–37](#)
Visual Studio, [1, 5](#)

W, X, Y, Z

Weak shared pointer
 (`std::weak_ptr`), [171–172](#)